

CLIENT SERVER TECHNOLOGY

C-132

Self Learning Material



Directorate of Distance Education

**SWAMI VIVEKANAND SUBHARTI UNIVERSITY
MEERUT-250005
UTTAR PRADESH**

SIM Module Developed by : M. K. Sharma

Reviewed by the Study Material Assessment Committee Comprising:

- 1. Lt. (Gen.) B.S. Rathore, Vice-Chancellor**
- 2. Dr. Sushmita Saxena, Pro-Vice-Chancellor**
- 3. Dr. Mohan Gupta**
- 4. Mr. Sumit Bhatnagar**
- 5. Mr. Rajeev Kumar**

Copyright © Laxmi Publications Pvt Ltd

No part of this publication which is material protected by this copyright notice may be reproduced or transmitted or utilized or stored in any form or by any means now known or hereinafter invented, electronic, digital or mechanical, including photocopying, scanning, recording or by any information storage or retrieval system, without prior permission from the publisher.

Information contained in this book has been published by Laxmi Publications Pvt Ltd and has been obtained by its authors from sources believed to be reliable and are correct to the best of their knowledge. However, the publisher and its author shall in no event be liable for any errors, omissions or damages arising out of use of this information and specially disclaim and implied warranties or merchantability or fitness for any particular use.

Published by : Laxmi Publications Pvt Ltd., 113, Golden House, Daryaganj, New Delhi-110 002.

Tel: 43532500, E-mail: info@laxmipublications.com

DEM-2230-48.90-CLIENT SERVER TECH C-132

Typeset at: Excellent Graphics, Delhi

Edition : 2017.

C- 01736/03/19

Printed at: Ajit Printing Press, Delhi

CONTENTS

Units	Page No.
1. Introduction to Client Server Technology	1
2. Fundamentals of Client Server Design	15
3. Transition to Client-Server Programming	32
4. Client Server Scheduling and Memory Management	58
5. Client Server Applications Design and Development	78

SYLLABUS

CLIENT SERVER TECHNOLOGY (BCA)

C-132

Unit - I

Client-Server Technology and its uses, historical development, client-server technology and heterogeneous computing, Distributed Computer, Computing platform forms.

Unit - II

Microprocessor integration and client server computing, implementations and scalability. Fundamentals of client server design, division of labor.

Unit - III

Transition to client-server programming; Interaction of client and server communication Techniques and protocols, implementing client server applications, multitasking with process and threads.

Unit - IV

Scheduling implementations, scheduler internals, preemptive Vs non-preemptive systems; synchronization-understanding and using semaphores, semaphore implementation in Novell Netware, windows NT and UNIX, Memory-management, Allocation, sharing and manipulating,

Unit - V

Client server computing with ORACLE-Overview of DBMS, client server relationships, ORACLE and client server computing, using SQL with SQL, *DBS, the ORACLE tools and design aids, SQL windows & Power Builder.

CHAPTER

1

INTRODUCTION TO CLIENT SERVER TECHNOLOGY

STRUCTURE

- 1.1 Client-Server Technology and Its Uses
- 1.2 Historical Development
- 1.3 Client-Server Technology and Heterogeneous Computing
- 1.4 Distributed Computer and Computing Platforms
 - *Summary*
 - *Review Questions*

1.1 CLIENT-SERVER TECHNOLOGY AND ITS USES

Having got a overview or a glimpse of what a client server computing is let us see what are the advantages of using the client server computing technology or in other words let us see the reasons why the client server computing has gained popularity among database world. Before knowing about the advantages and disadvantages of any technology it is good to have knowledge of the same. The below paragraph talk about what a Client Server computing is.

The Client/Server computing has three components in it namely:

- Client
- Server
- Network

Let us see the functions of each.

NOTES

Server: This is one that manages data or information in database. It also takes the role of coordination or assigning process when multiple clients functions which is very common in client/server computing technology.

Client: This is the application or in other words called as front end which the user uses to interact. This is the interface through which the user provides data to database or retrieve data from database.

Network: From the above we have seen that client and server are independent. But a link must be there to connect client and server for the application to function as a whole. This is done by network. This acts as a communication media between client and server.




Fig. 1.1 : Basic Client server model

Advantages of using the Client Server Technology

Cost Effective: From the above it is clear that each component namely the client and server in client server computing acts as a separate entity which can be connected by network. One of the advantages of having this type of design is it is possible for one to choose system as per the needs of the each component of a client server computing environment. In other words the system need of system component in client server computing is different from that of client component. Let us see this in brief. We know that server in client/server computing manages data in database. So for database management the server component system in client server computing must be a system with more memory and also high processor speed. One has to choose a system for a server in client server computing satisfying the above need. On the other hand the system for client component in a client server computing need not have the above design. That is system representing a client component in client server computing need not have more memory or high processor speed. But since the client represents user interface or front end application the system must be able to support graphical utilities or applications and also have or support secondary storage as per the needs. So when choosing system for a client component of a client server application one has to take this point into consideration and make the choice. Thus from the above explanation it is clear that the system chosen for server can be done separately and for client separately as per the needs. This helps the organization to make cost effective system by purchasing systems as per the needs only.

NOTES

Increased Performance: Performance is also increased without much effort. For instance if a organization decides to go for a system for server side to increase performance this can be done without affecting the user application because client component of client server computing need not be disturbed in this case.

Ease of effort and maintenance: Since the client and server component namely the user interface and database management are separated in the design of client server computing it eases the effort of database administrator and also developers. This is because developers can concentrate on user application alone. Also for instance if a user interface change occurs the system can be changed with little effort by making changes only in the client side without making server component changes.

Disadvantages

The above are some of the main advantages in client server computing. Though the advantages are numerous there are some disadvantage also in client server computing which made other technology like mainframes to take its shape and popularity. Some of the disadvantages are:

We have seen above that the client server computing is cost effective. It is true as said above and also because when compared with mainframes the cost of hardware is very less in client server computing. But in a client server computing it is not that hardware only is present but there are also other costs included like number of multiple clients' maintenance which includes users, developers, administrators to maintain database in server side and so on. When we consider all this the cost in client server computing is higher.

We know that mainframe is a centralized database managed system. Though the system is costlier in hardware point of view it provides reliable and secured system which is very much needed for big real time applications. This is because in these real times huge applications the main requirement is system must not stop and must run throughout without failure which is achieved by centralized maintained mainframe systems. But in client server computing it is not maintained centrally and the reliability is not much which means if the system got stopped in between then real time huge application would face huge loss in term of productivity and cost.

To truly understand how much of the Internet operates, including the Web, it is important to understand the concept of client server computing. The client server model is a form of distributed computing where one program (the client) communicates with another program (the server) for the purpose of exchanging information.

The client's responsibility is usually to:

1. Handle the user interface.
2. Translate the user's request into the desired protocol.
3. Send the request to the server.

4. Wait for the server's response.
5. Translate the response into "human-readable" results.
6. Present the results to the user.

NOTES

The server's functions include:

1. Listen for a client's query.
2. Process that query.
3. Return the results back to the client.

A typical client/server interaction goes like this:

1. The user runs client software to create a query.
2. The client connects to the server.
3. The client sends the query to the server.
4. The server analyzes the query.
5. The server computes the results of the query.
6. The server sends the results to the client.
7. The client presents the results to the user.
8. Repeat as necessary.

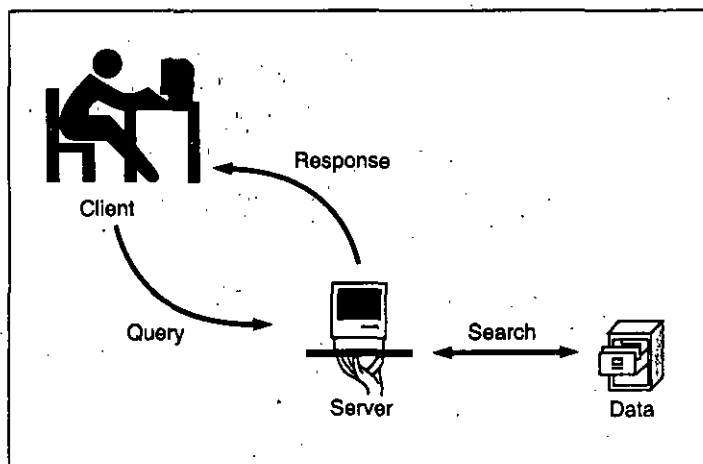


Fig. 1.2 : A typical client server interaction

This client server interaction is a lot like going to a French restaurant. At the restaurant, you (the user) are presented with a menu of choices by the waiter (the client). After making your selections, the waiter takes note of your choices, translates them into French, and presents them to the French chef (the server) in the kitchen. After the chef prepares your meal, the waiter returns with your diner (the results). Hopefully, the waiter returns with the items you selected, but not always; sometimes things get "lost in the translation."

Flexible user interface development is the most obvious advantage of client server computing. It is possible to create an interface that is independent of the server hosting the data. Therefore, the user interface of a client

NOTES

server application can be written on a Macintosh and the server can be written on a mainframe. Clients could be also written for DOS- or UNIX-based computers. This allows information to be stored in a central server and disseminated to different types of remote computers. Since the user interface is the responsibility of the client, the server has more computing resources to spend on analyzing queries and disseminating information. This is another major advantage of client server computing; it tends to use the strengths of divergent computing platforms to create more powerful applications. Although its computing and storage capabilities are dwarfed by those of the mainframe, there is no reason why a Macintosh could not be used as a server for less demanding applications.

In short, client server computing provides a mechanism for disparate computers to cooperate on a single computing task.

Uses of Client Server Technology

The **client-server model** of computing is a distributed application structure that partitions tasks or workloads between the providers of a resource or service, called servers, and service requesters, called clients. Often clients and servers communicate over a computer network on separate hardware, but both client and server may reside in the same system. A server machine is a host that is running one or more server programs which share their resources with clients. A client does not share any of its resources, but requests a server's content or service function. Clients therefore initiate communication sessions with servers which await incoming requests.

The *client-server* characteristic describes the relationship of cooperating programs in an application. The server component provides a function or service to one or many clients, which initiate requests for such services.

Functions such as email exchange, web access and database access, are built on the client-server model. Users accessing banking services from their computer use a web browser client to send a request to a web server at a bank. That program may in turn forward the request to its own database client program that sends a request to a database server at another bank computer to retrieve the account information. The balance is returned to the bank database client, which in turn serves it back to the web browser client displaying the results to the user. The client-server model has become one of the central ideas of network computing. Many business applications being written today use the client-server model. So do the Internet's main application protocols, such as HTTP, SMTP, Telnet, and DNS.

1.2 HISTORICAL DEVELOPMENT

The rapid proliferation of PCs in the workplace quickly exposed a number of their weaknesses. A stand-alone PC can be extremely inefficient. Any

NOTES


computing device requires some form of I/O system. The basic keyboard and monitor system is dedicated to one user, as it is hardly expected that two or more users will share the same PC at the same time. The same is not true for more specialized I/O devices, with which, for example, two or three printers attached to a mainframe or minicomputer environment can be accessed by any system user who can download a file to the printer. It might mean a walk to the I/O window to pick up a printout, but a few printers can meet many users' needs. In the stand-alone PC world, the printer is accessible only from the computer to which it is attached. Because the computer is a single-user system accessible only from the local keyboard, the printer cannot be shared, and therefore, must be purchased for each PC that needs to print; otherwise, PCs with dedicated printers must be available for anyone's use, in which case a user would take a file or data (usually on a floppy disk) to the printer station to print it.

This is affectionately referred to as sneakernet. It doesn't take a rocket scientist to note the waste in time, resources, and flexibility of this approach. We use printers here as just one example of the inefficiency of stand-alone PCs. Any specialized I/O device faces the same problems (i.e., plotters, scanners, and so on), along with such necessities as hard disk space (secondary storage) and even the processing power itself. Software is another resource that cannot be shared in stand-alone systems. Separate programs must be purchased and loaded on each station. If a department or company maintains database information, the database needs to be copied to any station that needs it. This is a sure-fire formula for inconsistent information or for creating an information bottleneck at some central administrative site. Finally, the stand-alone PC is isolated from the resources of the mainframe or minicomputer environment. Important information assets are not available, usually leading to two or more separate devices on each desk (such as a PC and a terminal into the corporate network). A vast array of computing power develops that is completely out of the control of the Information Technology (IT) group. The result can be (and often is) chaotic.

It rapidly became evident that a scheme was necessary to provide some level of interconnection. Enter the local area network (LAN). The LAN became the medium to connect these PCs to shared resources. However, the simple connection of resources and computers was not all that was required. Sharing these resources effectively requires a server. As an example of server function, consider again the problem of sharing a printer among a collection of end users. A printer is inherently a serial device (it prints one thing at a time). A printer cannot print a few characters submitted by user A, then a few from user B, and so on; it must print user A's complete document before printing user B's job. Simply connecting a printer to the network will not accomplish the serialization of printing, since users A and B are not synchronized with respect to job submission. A simple solution

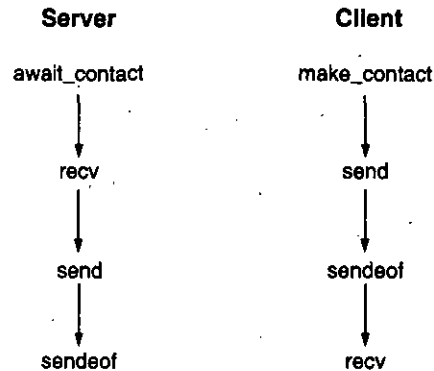
to this problem is to attach the printer to the network via a specialized processor, called a printer server. This processor accumulates inputs from users A and B, assembles each collection of inputs into a print job, and then sends the jobs to the printer in serial fashion.

NOTES



The printer server can also perform such tasks as initializing the printer and downloading fonts. The server must have substantial memory capability to assemble the various jobs, and it must contain the logic required to build a number of print queues (to prioritize the stream of printer jobs). A second example of a server's function involves a shared database connected to the network. In most systems, different users have different privileges associated with database access. Some might not be allowed access to certain files, others might be allowed to access these files to read information but not write to the files, while still others might have full read/write access. When multiple users can update files, a gate-keeping task must be performed, so that when user A has accessed a given file, user B is denied access to the file until user A is finished. Otherwise, user B could update a file at the central location while user A is still working on it, causing file overwrites. Some authority must perform file locking to assure that databases are correctly updated. In sophisticated systems, locking could be performed on a record (rather than a file) basis—user B can access any record that user A has not downloaded, but B cannot obtain access to a record currently being updated.

NOTES



The job of the file (or database) server is to enforce security measures and guarantee consistency of the shared database. The file server must have substantial resources to store all the requisite databases and enough processing power to respond quickly to the many requests submitted via the network. Many other server types are available. For example, a communications server might manage requests for access to remote resources (offsite resources that are not maintained locally). This server would allow all requests for remote communication to be funneled through a common processor, and it would provide an attachment point for a link to a WAN. Application servers might perform specialized computational tasks (graphics is a good example), minimizing the requirement for sophisticated hardware deployed at every network location. Servers are sometimes simply PCs, but they are often specialized high-speed machines called workstations. In some environments, the servers might even be minicomputers or mainframes. Those computers that do not provide a server function are typically called clients, and most PC networks are client/server oriented.

1.3 CLIENT-SERVER TECHNOLOGY AND HETEROGENEOUS COMPUTING

A client server model is used to describe the generic services provided in a heterogeneous computing environment. While other models may also apply, client server models serve well as a means of describing the services outlined in this section.

In a client server model, a node in a heterogeneous computing environment (e.g., personal computer, workstation, minicomputer, mainframe) may be characterized either as a client, a server, or a client server. A *client* node provides no services to any other node. A client is only able to be the recipient of services provided by other nodes. A node that provides services is a *server*. Often a server can also be a client. Such a node is referred to as a *client server*. The term *client* may also refer to a process which runs on a client node or to the person on whose behalf a client process is acting. Similarly, the term *server* may also refer to a process which runs on a server node. In this report, the term *client* refers to a *client node* and the term *server* refers to a *server node*.

NOTES

In next figure, a personal computer and a workstation are clients. They provide no services to any other node, but they may be the recipient of a service, such as mounting a file system or directory, from either the server or client server. The client server may also be the recipient of services.




Fig. 1.3 : Client/Server Model

Originally, nodes of computer networks were exclusively mainframes and minicomputers. The user accessed the network only by means of a terminal connected to one of the nodes. The services provided by the network were usually login (sometimes called *virtual terminal* or *remote login*), file transfer, and mail. Each node of the network was a client/server for these three services. Users were provided with an interactive terminal session with any node on the network. They were able to transfer files between any two nodes. Finally, they were able to send mail to any node and receive mail from any node.

With the advent of the personal computer, networks were implemented whose nodes were exclusively small computers. Such a *personal computer* network provided file service and mail. Each node of a personal computer network was usually either a client or a server and typically, there were no client/server nodes. There was no login service provided since each personal computer ran applications locally. Although applications may be downloaded from a server, their primary function was to provide file service for data (*i.e.*, transparent access to data on mass storage). File transfer between clients was accomplished by copying files to and from a server. Additionally, servers sometimes acted as a post office for electronic mail. Clients accessed the mail server to send and receive mail.

This distinction between large computer networks and personal computer networks has somewhat disappeared. In a heterogeneous computing environment, the services of the large computer network and the personal computer network are integrated into three basic types of services which are available to all sizes of clients from all sizes of servers or client servers. In a heterogeneous computing environment, clients, servers, and client servers may be, for

example, mainframes, minicomputers, workstations, or personal computers. Most services can be grouped into three major categories: distributed file system, distributed computing, and messaging.

NOTES

Distributed File System

The concept of a *distributed file system* is the generalization of the concept of file service. File service provides a client with transparent access to part of the mass storage of a remote server. Transparent access means that, to the user on a client, the remote mass storage of the server is available on the client as though it were mass storage locally connected to the client.

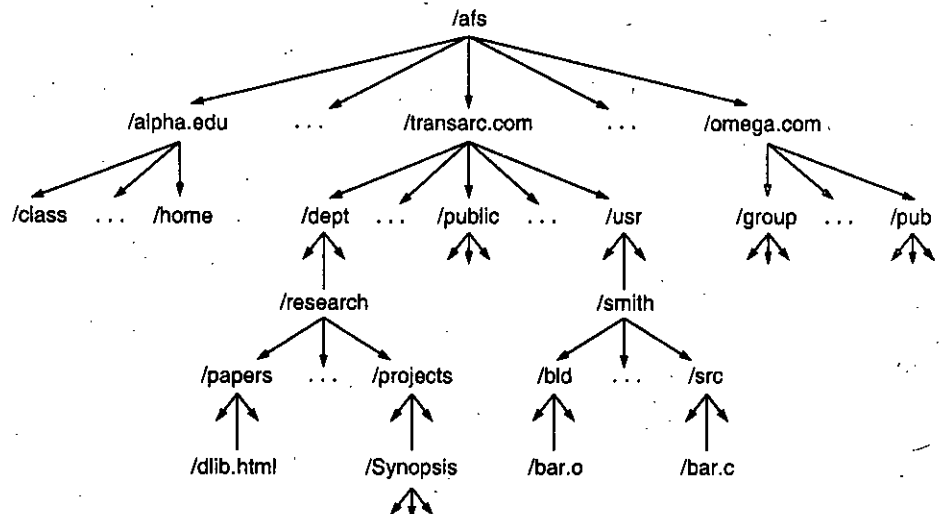


Fig. 1.4 : An example of Distributed file system

Figure 1.4 illustrates the concept of transparent access. The "/" notation is used to indicate the location of a file within the directory tree of the client. The "." and "." notation is used to indicate the location of a file on the server (see sec.). The client has, as part of its file system, a part of the file system of a remote system which is providing file service. The directory tree which begins at /users/Jane on the clients directory tree physically resides on the server's directory tree at the point UserDisk:[NetUsers.Jane] which is part of the server's mass storage device UserDisk. Note that, in this example, although the syntax of file reference differs between the client and the server, the client uses its file reference syntax (i.e., the "/" notation) for all file references including those which physically reside on the server. The fact that the client and server use different file reference syntax notations implies that the client and server have different operating systems. In a heterogeneous computing environment, a client may have access to several servers simultaneously. In addition, on some servers, a client may have transparent access to peripheral devices on servers other than mass storage, such as, printers and modems.

NOTES

Since a client and server may be from different producers, a user on a client is provided with transparent access to remote peripherals in the user interface paradigm of the client, not the server. This means that, once remote mass storage is attached, users issue the same commands (e.g., copy, rename, delete) to access remote mass storage as they do to access local mass storage. Similarly, once a remote printer is attached, the commands to print on the remote printer are the same as the commands to print on the local printer. The file system/device model (i.e., how file or device references are supported in the command language) used by each client is the access provided by servers to clients in a heterogeneous computing environment.

1.4 DISTRIBUTED COMPUTER AND COMPUTING PLATFORMS

The second major category of service in a heterogeneous computing environment is *distributed computing*. Distributed computing refers to the concept of running an application or applications on remote node(s). Distributed computing is not the downloading of an application to be run locally. Such downloading is a service of the distributed file system. Distributed computing is the execution of an application on a remote server or servers with the results transmitted over the network to the client. Distributed computing takes four forms: login, remote execution, transaction processing, and cooperative processing.




Fig. 1.5 : Login (Client A) vs. Remote Execution (Client B)

Login

Login is the service which provides a user on a client with an interactive session on a server. By means of the login service, the client behaves as though it were a terminal directly connected to a host where the host is a login server on the network. Typically, the login session is conducted using the user interface of the login server. Some clients are able to have several login sessions to multiple servers active simultaneously. Normally,

NOTES

each login session requires a dedicated process on the client and a dedicated process on the server. For those clients with mouse/window user interfaces, each login session is usually a separate window. If the client is also a server, then remote interactive applications can access local mass storage for data by means of the distributed file system. If not, then file transfer to a server may be used to allow access to data. Login permits access to remote applications which are highly interactive in nature.

Figure 1.5 shows the interaction between a user on client **A** and a login server in the typical login session. The purpose of the session is to obtain a directory listing of a directory on the login server. The **bold** indicates what the user enters.

In the future, the need for a login service will diminish. Users will interact with applications as though the application were resident locally, whereas, in fact, the processing of the application may be taking place on several different nodes.

Remote Execution

Remote execution is the service which provides a user on a client access to remote applications which are not highly interactive in nature. Remote execution is differentiated from login in several ways. With login, the user on the client interacts with each server in the command language (*i.e.*, user interface paradigm) of the server. With remote execution, the user issues commands to the server using the user interface of the client, *i.e.*, commands executed by the server are issued in the same language as commands to the client. From an implementation point of view, a login session requires a process on the client and a process on the server dedicated to each interactive session throughout its establishment. On the other hand, remote execution normally only requires the client to execute a single process in order to access multiple servers at the same time. Moreover, a single process on the server is usually sufficient to manage processing requests from all clients.

Figure 1.5 shows the interaction between a user on client **B** and a remote execution server. As the case with the user on client **A** accessing a login server, the user on client **B** is obtaining a directory listing from a remote system. The **bold** indicates what the user enters. Note that the user on client **B** is not *logged into* the server, but is simply entering a command which happens to be carried out on the server. In many cases, remote execution happens without the user's knowledge. With remote execution, all operations can appear as though they are happening locally even when they take place remotely.

Remote execution is also different from the venerable concept of *remote job entry*. Remote job entry is the remote initiation of batch processing on a host. Remote execution on a server normally proceeds at interactive priorities (*i.e.*, the same priorities as interactive sessions). Batch priorities are usually less than interactive priorities. In addition, the concept of

remote execution includes the concept of applications on different nodes closely interact with each other in a highly synchronized manner.

Transaction Processing

The concept of *transaction processing* services evolved from the requirements of database applications. A *transaction* is an update (i.e., change) made to information in a database. Many transactions involve a query as well. The database is queried, a decision made on the basis of the information obtained from the query, and then a change made to the database. An application maintaining a database on a server in a heterogeneous computing environment must be capable of managing transactions, which arrive virtually simultaneously from clients, in a manner which insures the integrity of the information.

In a heterogeneous computing environment, transaction processing services are provided in order to meet the special needs of distributed database applications. A transaction applied against a database on a server can cause the server to generate transactions against databases on other servers, and so on. The initial transaction generates a flow of transactions which may be depicted as a flow through a tree structure. Each leaf of the tree represents a server running an application maintaining a database. Transaction processing services help the server database application control and synchronize these transaction flows so that the information contained in all of the databases is reliable and up-to-date.

Cooperative Processing

The availability of remote execution services and transaction processing services has given rise to the concept of a *distributed application*. A distributed application is an application in which the interaction between the user and the application appears as though all of the processing is taking place locally, but, in fact, the processing may be distributed across several nodes of the network. The most common example of a distributed application can be found in a distributed database application. The user on a client initiates queries to a database whose data is distributed among several nodes of the network. To the user, the data appears to be local. A server (or servers) accepts the queries, generates the information, and sends the results back to the client. Distributed applications make use of *cooperative processing* services. The application programmer uses cooperative processing services for processing required by the application and, in some cases, need not be aware of the location where the processing is taking place. The cooperative processing service may be able to determine automatically where processing is to take place. Cooperative processing services include the functionality provided by both remote execution and transaction processing services. In the future, the concepts of cooperative processing and distributed applications should replace the concepts of login, remote execution, and transaction processing. Users will only interact with distributed applications that are implemented using cooperative processing services.

NOTES

NOTES

SUMMARY

1. Server manages data or information in database. It also takes the role of coordination or assigning process when multiple clients functions which is very common in client/server computing technology.
2. It is possible for one to choose system as per the needs of the each component of a client server computing environment.
3. Performance is also increased without much effort.
4. Since the client and server component namely the user interface and database management are separated in the design of client server computing it eases the effort of database administrator and also developers.
5. Flexible user interface development is the most obvious advantage of client server computing. It is possible to create an interface that is independent of the server hosting the data.
6. The client server model of computing is a distributed application structure that partitions tasks or workloads between the providers of a resource or service, called servers, and service requesters, called clients.
7. A client server model is used to describe the generic services provided in a heterogeneous computing environment.
8. Login is the service which provides a user on a client with an interactive session on a server.

REVIEW QUESTIONS

1. How you can define the term client server computing?
2. Why we need client server model for data base applications?
3. Explain the entity involve to design a client server model.
4. What are the main application areas of client server model?
5. Write a brief history of client server computing.
6. Do you think in the age of online application, we can still use client server model?
7. Write the advantages and disadvantages of client server model.
8. Design a client server model for a small local bank of your area.
9. Find out the software and operating systems that support a client server model.
10. Explain Distributed Computer and Computing platforms.

CHAPTER

2

FUNDAMENTALS OF CLIENT SERVER DESIGN

STRUCTURE

- 2.1 Microprocessor integration and client server computing
- 2.2 Implementations and scalability
- 2.3 Fundamentals of client server design
- 2.4 Division of labor
 - *Summary*
 - *Review Questions*

2.1 MICROPROCESSOR INTEGRATION AND CLIENT SERVER COMPUTING

Modern client-server computing requires communication between client and server over heterogeneous networks and operating environments. Communication needs have gone beyond SQL statements to procedure transactions and object-traffic. Middleware software has evolved to serve this communications need, but solutions have been extremely complex, costly, and proprietary. The World Wide Web, the Common Gateway Interface and the Java environment, have become powerful alternatives to traditional middleware. Approaches to join Web and Distributed Object technologies are being developed.

In April of 1995 a review of client/server computer was published in BYTE magazine. The authors described a dazzling range of technologies that would take computing into the "intergalactic" realm. Nowhere is the word "web" or "Internet" mentioned.

NOTES

Precsely one year later an IBM Executive Overview, published on IBM's website, declared that "the web changes the client-server game" (Furey, 1996). In the same month BYTE declared "the Internet and the World Wide Web may be all the middleware you need". Everywhere, Sun's Java programming language is widely proclaimed to be the ultimate solution to all client-server dilemmas.

Client-server computing is roughly as old as the Internet; both came into their own in the early 1980s. At the time most business computing was done using mainframe systems. Mainframes centralized data and processes, sending information back to terminals for viewing. The terminals could interpret formatting commands and display data, but they had little processing power.

Microprocessors changed the economics of the mainframe world they made it possible to create workstations or personal computers that could do a great deal more than interpret formatting codes. Database and hardware vendors challenged the mainframe world with client-server solutions. Client software, running on a local computer, allowed a much greater range of local data representation and processing, and supported graphical interfaces. Server software ran on more powerful microprocessor systems, and responding to client queries with streams of data. Systems could be created by small groups without the need for centralized mainframes. SQL database management systems (DBMS), from vendors such as Oracle and Sybase, became widely used.

SQL database servers have been successful, but multivendor interoperability has been limited by proprietary solutions for transaction support and stored procedures. In the 1990s many complex client-server implementations have foundered. The costs of operating and maintaining a network of personal computers has exceeded traditional mainframe costs (Furey, 1996). Users are overwhelmed by complex architectures for next-generation client/server systems featuring transaction processing monitors, multimedia document repositories, and intricate standards for object interoperability.

Integration with Distributed Computing

Distributed computing is the term used for implementation of technologies across heterogeneous environments. For operating systems, heterogeneous computing means the ability to communicate with other operating systems and protocols. Distributed computing is a complex architecture. It involves rearchitecture of applications, redevelopment of systems and increased efficiency in maintaining a network as a whole. Many distributed nodes work on behalf of one requesting client. This makes the system fault tolerant and decentralized, which is an obvious advantage over centralized systems. For the technology to become effective and revolutionary, developers of distributed applications have to do everything possible to minimize the complexity of development and maintenance and integrate their software

with disparate platforms. Client/Server application designing necessitates the modularization of applications and their functions into discrete components. These components must be bounded only by encapsulated data and functions that may be moved between the systems. This design model gives Client/Server software more adaptability and flexibility.

NOTES

2.2 IMPLEMENTATIONS AND SCALABILITY

Scalability: The modular nature of the Client Server system may be replaced without adversely affecting the rest of the system. For example, it is possible to upgrade the server to a more powerful machine with no visible changes to the end user. This ability to change component system makes Client Server systems especially receptive to new technologies in both hardware and software.

2.3 FUNDAMENTALS OF CLIENT SERVER DESIGN

The term Client Server was first used in the 1980s in reference to personal computers (PCs) on a network. The actual Client Server model started gaining acceptance in the late 1980s. The term Client/Server is in reality a logical concept. The client and server components may not exist on distinct physical hardware. A single machine can be both a client and a server depending on the software configuration. The Client Server technology is a model, for the interaction between simultaneously executing software processes. The term architecture refers to the logical structure and functional characteristics of a system, including the way they interact with each other in terms of computer hardware, software and the links between them.

In case of Client Server systems, the architecture means the way clients and servers along with the requisite software are configured with each others. Client Server architecture is based on the hardware and the software components that interact to form a system. The limitations of file sharing architectures led to the emergence of the Client Server architecture. This approach introduced a database server to replace the file server. Using a Relational Database Management System (RDBMS), user queries could be answered directly. The Client Server architecture reduced network traffic by providing a query response rather than total file transfer. It improves multi-user updating through a GUI front-end to a shared database. In Client Server architectures, Remote Procedure Calls (RPCs) or Structural Query Language (SQL) statements are typically used to communicate between the client and server.

Broadly, there are three types of Client Server systems in existence.

(i) Two-tier

(ii) Three-tier

(iii) N-Tier

NOTES

(i) Two-tier Client Server Model

The application processing is done separately for database queries and updates and for business logic processing and user interface presentation. Usually, the network binds the back-end of an application to the front-end, although both tiers can be present on the same hardware.

Sometimes, the application logic (the real business logic) is located in both the client program and in the database itself. Quiet often, the business logic is merged into the presentation logic on the client side. As a result, code maintenance and reusability become difficult to achieve on the client side. On the database side, logic is often developed using stored procedures.

In the two-tier architecture, if the Client/Server application has a number of business rules needed to be processed, then those rules can reside at either the Client or at the Server. The Fig. 2.1 below clarifies this situation.




Fig. 2.1 : 2-tier architecture

2-tier architecture is used to describe client/server systems where the client requests resources and the server responds directly to the request, using its own resources. This means that the server does not call on another application in order to provide part of the service.

There are several advantages of two-tier systems

- Availability of well-integrated PC-based tools like, Power Builder, MS Access, 4 GL tools provided by the RDBMS manufacturer, remote SQL, ODBC.

- Tools are relatively inexpensive.
- Least complicated to implement.
- PC-based tools show Rapid Application Development (RAD) *i.e.*, the application can be developed in a comparatively short time.
- The 2-tier Client/Server provides much more attractive graphical user interface.

(GUI) applications than was possible with earlier technology.

Architecture maintains a persistent connection between the client and database, thereby eliminating overhead associated with the opening and closing of connections.

- Faster than three-tier implementation.
- Offers a great deal of flexibility and simplicity in management.

Conversely, a two-tier architecture has some disadvantages

As the application development is done on client side, maintenance cost of application, as well as client side tools etc. is expensive. That is why in 2-tier architecture the client is called 'fat client'.

- Increased network load: Since actual processing of data takes on the remote client, the data has to be transported over the network. This leads to the increased network stress.
- Applications are loaded on individual PC *i.e.* each application is bound to an individual PC. For this reason, the application logic cannot be reused.
- Due to dynamic business scenario, business processes/logic have to be changed.

These changed processes have to be implemented in all individual PCs. Not only that, the programs have to undergo quality control to check whether all the programs generate the same result or not.

- Software distribution procedure is complicated in 2-tier Client/Server model. As all the application logic is executed on the PCs, all these machine have to be updated in case of a new release. The procedure is complicated, expensive, prone to errors and time consuming.
- PCs are considered to be weak in terms of security *i.e.*, they are relatively easy.
- To crack.
- Most currently available drivers require that native libraries be loaded on a client machine.
- Load configurations must be maintained for native code if required by the driver.

Problem areas are encountered upon implementing this architecture on the Internet.

NOTES

(ii) Three-Tier Client Server Model

In 3-tier architecture, there is an intermediary level, meaning the architecture is generally split up between:

NOTES

1. A client, *i.e.* the computer, which requests the resources, equipped with a user interface (usually a web browser) for presentation purposes.
2. The application server (also called **middleware**), whose task it is to provide the requested resources, but by calling on another server.
3. The data server, which provides the application server with the data it requires.

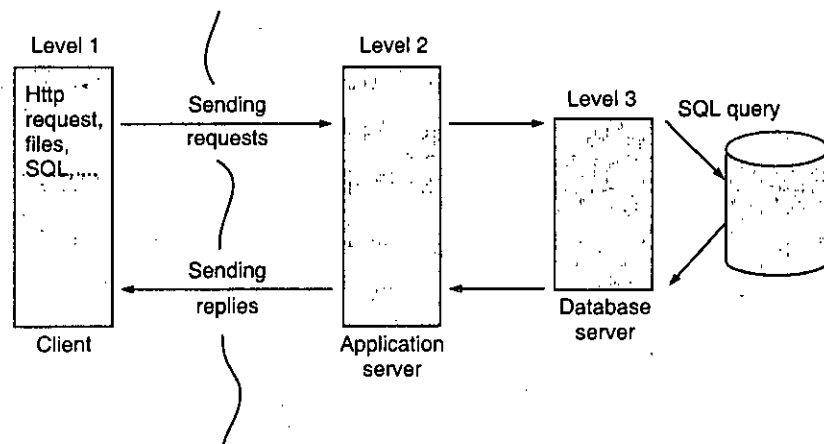


Fig. 2.2 : Three-tier architecture application

In three-tier architecture application responsibilities are divided into three logical categories (in other words, the business system must provide three types of main services).

Presentation (GUI) or user services: Include maintaining the graphical user interface and generating what users see on the monitor. Presentation Logic dealing with:

- Screen formatting.
- Windows management.
- Input editing.
- What-if analysis.

Application services or business rules: These include executing applications and controlling program flow. Business logic dealing with:

- Domain and range validation.
- Data dependency validation.
- Request/response architecture of Inter Process Communication level.

Database services or data server: Which refers to the management of underlying databases. Server logic deals with:

- Data access.

- Data management.
- Data security.
- SQL parsing.

Based on these three components, the three-tier architecture of Client Server system is shown in Fig. 2.2 In three-tier model, a third server is employed to handle requests from the client and then pass them off to the database server.

The boundaries between tiers are logical. One can run 3-tiers in one and the same machine. The important fact is that the system is neatly structured and well-planned definitions of the software boundaries exist between the different tiers. Some of the advantages of using three-tier model include:

- Application maintenance is centralized with the transfer of the business logic for many end users into a single application server. This eliminates the concern of software distribution that are problematic in the traditional two-tier Client Server model.
- Clear separation of user-interface-control and data presentation from application logic. Through this separation more clients are able to have access to a wide variety of server applications.

The two main advantages for client-applications are clear:

Quicker development through the reuse of pre-built business-logic components and a shorter test phase, because the server-components have already been tested.

- Many users are able to access a wide variety of server applications, as all application logic are loaded in the applications server.
- As a rule servers are "trusted" systems. Their authorization is simpler than that of thousands of "untrusted" client-PCs. Data protection and security is simpler to obtain. Therefore, it makes sense to run critical business processes that work with security sensitive data, on the server.
- Redefinition of the storage strategy won't influence the clients. RDBMS' offer a certain independence from storage details for the clients. However, cases like changing table attributes make it necessary to adapt the client's application.

In the future, even radical changes, like switching from an RDBMS to an OODBMS, won't influence the client. In well-designed systems, the client still accesses data over a stable and well-designed interface, which encapsulates all the storage details.

- Load balancing is easier with the separation of the core business logic from the database server.
- Dynamic load balancing: if bottlenecks in terms of performance occur, the server process can be moved to other servers at runtime.

NOTES

NOTES

- Business objects and data storage should be brought as close together as possible. Ideally, they should be together physically on the same server. This way network load for complex access can be reduced.
- The need for less expensive hardware because the client is 'thin'.
- Change management is easier and faster to execute. This is because a component/program logic/business logic is implemented on the server rather than furnishing

Comparing Both Types of Architecture

2-tier architecture is therefore a client-server architecture where the server is versatile, *i.e.* it is capable of directly responding to all of the client's resource requests.

In 3-tier architecture however, the server-level applications are remote from one another, *i.e.* each server is specialized with a certain task (for example: web server/database server). 3-tier architecture provides:

- A greater degree of flexibility.
- Increased security, as security can be defined for each service, and at each level.
- Increased performance, as tasks are shared between servers.

(iii) Multi-Tiered or N-tier Client Server Model

In 3-tier architecture, each server (tier 2 and 3) performs a specialized task (a service). A server can therefore use services from other servers in order to provide its own service. As a result, 3-tier architecture is potentially an n-tiered architecture.




Fig. 2.3 : Multi-Tiered or N-tier Client Server Model

Components of Client Server Model

Client Server architecture is based on hardware and software components that interact to client form a server system. The system includes mainly three components.

- (i) Hardware (client and server).
- (ii) Software (which make hardware operational).
- (iii) Communication middleware. (associated with a network which are used to link the hardware and software).

The client is any computer **process** that requests services from server. The client uses the services provided by one or more server processors. The client is also known as the **front-end application**, reflecting that the end user usually interacts with the client process.

The server is any computer **process** providing the services to the client and also supports multiple and simultaneous clients requests. The server is also known as **back-end application**, reflecting the fact that the server process provides the background services for the client process.

The communication middleware is any computer **process through** which client and server communicate. Middleware is used to integrate application programs and other software components in a distributed environment. Also known as **communication layer**. Communication layer is made up of several layers of software that aids the transmission of data and control information between Client and Server. Communication middleware is usually associated with a network. The Fig.2.4 below gives a general structure of Client Server System.




Fig. 2.4 : Structure of a Client Server System

Now as the definition reveals, clients are treated as the front-end application and the server as the back-end application, the Fig. 2.5 shows the front-end and back-end functionality.




Fig. 2.5 : Front-end and Back-end Functionality

NOTES

Principles Behind Client Server Systems

NOTES

The components of the Client Server architecture must conform to some basic principles, if they are to interact properly. These principles must be uniformly applicable to client, server, and to communication middleware components. Generally, these principles generating the Client Server architecture constitute the foundation on which most current generation Client Server system are built. Some of the main principles are as follows:

- (i) Hardware independence.
- (ii) Software independence.
- (iii) Open access to services.
- (iv) Process distribution.
- (v) Standards.

(i) Hardware independence: The principles of hardware independence requires that the Client, Server, and communication middleware, processes run on multiple hardware platforms (IBM, DEC, Compaq, Apple, and so on) without any functional differences.

(ii) Software independence: The principles of software independence requires that the Client, Server, and communication middleware processes support multiple operating systems (such as Windows 98, Windows NT, Apple Mac system, OS/2, Linux, and Unix) multiple network protocols (such as IPX, and TCP/IP), and multiple application (spreadsheet, database electronic mail and so on).

(iii) Open access to services: All client in the system must have open (unrestricted) access to all the services provided within the network, and these services must not be dependent on the location of the client or the server. A key issue is that the services should be provided on demand to the client. In fact, the provision of on demand service is one of the main objectives of Client/Server computing model.

(iv) Process distribution: A primary identifying characteristic of Client Server system is that the processing of information is distributed among Clients and Servers. The division of the application-processing load must conform to the following rules:

- Client and server processes must be autonomous entities with clearly defined boundaries and functions. This property enables us to clearly define the functionality of each side, and it enhances the modularity and flexibility of the system.
- Local utilization of resources (at both client and server sides) must be maximized.

The client and server process must fully utilize the processing power of the host computers. This property enables the system to assign functionality to the computer best suited to the task. In other words, to best utilize all resources, the server process must be shared among all client processes;

that is, a server process should service multiple requests from multiple clients.

- Scalability and flexibility requires that the client and server process be easily upgradeable to run on more powerful hardware and software platforms. This property extends the functionality of Client Server processes when they are called upon to provide the additional capabilities or better performance.
- Interoperability and integration requires that client and server processes be seamlessly integrated to form a system. Swapping a server process must be transparent the client process.

NOTES

CLIENT SIDE COMPONENTS

As we know, the client is any process that requests services from the server process. The client is proactive and will, therefore, always initiate the conversation with the server. The client includes the software and hardware components. The desirable client software and hardware feature are:

- (i) Powerful hardware.
- (ii) An operating system capable of multitasking.
- (iii) Communication capabilities.
- (iv) A graphical user interface (GUI).

(i) Powerful hardware: Because client processes typically requires a lot of hardware resources, they should be stationed on a computer with sufficient computing power, such as fast Pentium II, III, or RISC workstations. Such processing power system handles multiple data types, such as voice, image, video, and so on. Client processes also require large amount of hard disk space and physical memory, the more such a resource is available, the better.

(ii) An operating system capable of multitasking: The client should have access to an operating system with at least some multitasking capabilities. Microsoft Windows 98 and XP are currently the most common client platforms. Windows 98 and XP provide access to memory, pre-emptive multitasking capabilities, and a graphical user interface, which makes windows the platform of choice in a majority of Client Server implementations. However, Windows NT, Windows 2000 server, OS/2 from IBM corporation, and the many "flavours" of UNIX, including Linux are well-suited to handle the Client/Server processing that is largely done at the server side of the Client/Server equation.

(iii) Communication capabilities: To interact efficiently in a Client Server environment, the client computer must be able to connect and communicate with the other components in a network environment. Therefore, the combination of hardware and operating system must also provide adequate connectivity to multiple network operating systems. The reason for requiring

NOTES

a client computer to be capable of connecting and accessing multiple network operating systems is simple services may be located in different networks.

(iv) A graphical user interface (GUI): The client application, or front-end, runs on top of the operating system and connects with the communication middleware to access services available in the network. Several third generation programming languages (3GLs) and fourth generation languages (4GLs) can be used to create the front-end application. Most front-end applications are GUI-based to hide the complexity of the Client/Server components from the end user.

SERVER SIDE COMPONENTS

As we have already discussed, the server is any process that provides services to the client process. The server is active because it always waits for the client's request. The services provided by server are:

(i) File services: For a LAN environment in which a computer with a big, fast hard disk is shared among different users, a client connected to the network can store files on the file server as if it were another local hard disk.

(ii) Print services: For a LAN environment in which a PC with one or more printers attached is shared among several clients, a client can access any one of the printers as if it were directly attached to its own computer. The data to be printed travel from the client's PC to the server printer PC where they are temporarily stored on the hard disk. When the client finishes the printing job, the data is moved from the hard disk on the print server to the appropriate printer.

(iii) Fax services: This requires at least one server equipped (internally or externally) with a fax device. The client PC need not have a fax or even a phone line connection. Instead, the client submits the data to be faxed to the fax server with the required information; such as the fax number or name of the receiver. The fax server will schedule the fax, dial the fax number, and transmit the fax. The fax server should also be able to handle any problems derived from the process.

(iv) Communication services: That let the client PCs connected to the communications server access other host computers or services to which the client is not directly connected. For example, a communication server allows a client PC to dial out to access board, a remote LA location, and so on.

(v) Database services: Which constitute the most common and most successful Client Server implementation. Given the existence of database server, the client sends SQL request to the server. The server receives the SQL code, validates it, executes it, and send only the result to the client. The data and the database engine are located in the database server computer.

(vi) Transaction services: Which are provided by transaction servers that are connected to the database server. A transaction server contains

NOTES

the database transaction code or procedures that manipulate the data in database. A front-end application in a client computer sends a request to the transaction server to execute a specific procedure store on the database server. No SQL code travels through the network. Transaction servers reduce network traffic and provide better performance than database servers.

(vii) Groupware services: Liable to store semi-structured information like Text, image, mail, bulletin boards, flow of work. Groupware Server provides services, which put people in contact with other people, that is because "groupware" is an ill-defined classification. Protocols differ from product to product. For examples: Lotus Notes/Domino, Microsoft Exchange.

(viii) Object application services: Communicating distributed objects reside on server. Object server provides access to those objects from client objects. Object Application Servers are responsible for Sharing distributed objects across the network. Object Application Servers uses the protocols that are usually some kind of Object Request Broker (ORB). Each distributed object can have one or more remote method. ORB locates an instance of the object server class, invokes the requested method, and returns the results to the client object. Object Application Server provides an ORB and application servers to implement this.

(ix) Web application services: Some documents, data, etc., reside on web servers. Web application provides access to documents and other data. "Thin" clients typically use a web browser to request those documents. Such services provide the sharing of the documents across intranets, or across the Internet (or extranets). The most commonly used protocol is HTTP (Hyper Text Transfer Protocol). Web application servers are now augmenting simple web servers.

(x) Miscellaneous services: These include CD-ROM, video card, backup, and so on. Like the client, the server also has hardware and software components. The hardware components include the computer, CPU, memory, hard disk, video card, network card, and so on. The computer that houses the server process should be the more powerful computer than the "average" client computer because the server process must be able to handle concurrent requests from multiple clients.

The server application, or back-end, runs on the top of the operating system and interacts with the communication middleware components to "listen" for the client request for the services. Unlike the front-end client processes, the server process need not be GUI based. Keep in mind that back-end application interacts with operating system (network or stand alone) to access local resources (hard disk, memory, CPU cycle, and so on). The back-end server constantly "listens" for client requests. Once a request is received the server processes it locally. The server knows how to process the request; the client tells the server only what it needs do not how to do it. When the request is met, the answer is sent back to the client through the communication middleware.

NOTES

The server hardware characteristics depend upon the extent of the required services. For example, a database is to be used in a network of fifty clients may require a computer with the following minimum characteristics:

- Fast CPU (RISC, Pentium, Power PC, or multiprocessor)
- Fault tolerant capabilities:
- Dual power supply to prevent power supply problem.
- Standby power supply to protect against power line failure.
- Expandability of CPU, memory disk, and peripherals.
- Bus support for multiple add-on boards.
- Multiple communication options.

In theory, any computer process that can be clearly divided into client and server components can be implemented through the Client/Server model. If properly implemented, the Client Server architectural principles for process distribution are translated into the following server process benefits:

- **Location independence.** The server process can be located anywhere in the network.
- **Resource optimization.** The server process may be shared.
- **Scalability.** The server process can be upgraded to run on more powerful platforms.
- **Interoperability and integration.** The server process should be able to work in a "Plug and Play" environment.

These benefits added to hardware and software independence principles of the Client Server computing model, facilitate the integration of PCs, minicomputer, and mainframes in a nearly seamless environment.

2.4 DIVISION OF LABOR

The tier architecture based client server model which is also known as client-server architecture, mainly focuses on the division of labor in an organization. The meaning is that we can divide a large task to many machines known as clients.

In computing terminology, both "client" and "server" refer to computers that are used for different purposes. A client is a small computer that accesses a server through a network. For example, in an organization, an employee logs in to the client machine to access the files and applications running on a server machine. This two-tier architecture is also known as client-server architecture which mainly focuses on the division of labor in an organization. A server machine is a large-capacity computer that can store a wide variety of files such as application and data files. There are various types of servers, such as; application server, file server, web server, database server, print server, proxy server, game server, standalone server, etc. A client can be classified into fat, thin, and hybrid. A fat client supports

both local storage and local processing. A thin client is a less powerful machine with minimum hardware installed. It usually utilizes the resources of a host machine and relies on the server to perform any data processing. The primary job of a thin client is just to graphically display the images provided by an application server. A hybrid client processes locally but relies on the server for data storage.

Some application servers may require users to log in from their client machines in order to access specific applications utilizing the client-server architecture. The client machines cannot only access the applications and data files, but they can also use the processor of the server to perform certain tasks without having to add any additional hardware resources to the client machine.

The client computer usually contains more end-user software than the server computer. A server usually contains more operating system components. Multiple users can log into a server at the same time. A client machine is simple and inexpensive whereas a server machine is more powerful and expensive.

The main difference between a client machine and a server machine is in its performance. The client machines are considered optimal for applications which require speedy start-up times. A server machine is considered optimal for applications where the emphasis is more on performance.




Fig. 2.6 : Division of labor using many clients

The term Client server refers to the concept of sharing the work involved in processing data between the client computer and the most

NOTES

powerful server computer. The client server network is the most efficient way to provide: Databases and management of applications such as Spreadsheets, Accounting, Communications and Document management.

NOTES

The client server model is basically an implementation of distributed or cooperative processing. At the heart of the model is the concept of splitting application functions between a client and a server processor. The division of labor between the different processors enables the application designer to place an application function on the processor that is most appropriate for that function. This lets the software designer optimize the use of processors—providing the greatest possible return on investment for the hardware. Client server application design also lets the application provider mask the actual location of application function. The user often does not know where a specific operation is executing. The entire function may execute in either the PC or server, or the function may be split between them. This masking of application function locations enables system implementers to upgrade portions of a system over time with a minimum disruption of application operations, while protecting the investment in existing hardware and software.

SUMMARY

1. Distributed computing is the term used for implementation of technologies across heterogeneous environments. For operating systems, heterogeneous computing means the ability to communicate with other operating systems and protocols. Distributed computing is a complex architecture.
2. 2-tier architecture is used to describe client/server systems where the client requests resources and the server responds directly to the request, using its own resources. This means that the server does not call on another application in order to provide part of the service.
3. Since actual processing of data takes on the remote client, the data has to be transported over the network. This leads to the increased network stress.
4. 2-tier architecture is therefore a client-server architecture where the server is versatile, i.e. it is capable of directly responding to all of the client's resource requests.
5. In 3-tier architecture however, the server-level applications are remote from one another, i.e. each server is specialized with a certain task (for example: web server/database server).
6. The communication middleware is any computer process through which client and server communicate.
7. A client is a small computer that accesses a server through a network.
8. A server machine is a large-capacity computer that can store a wide variety of files such as application and data files.

REVIEW QUESTIONS

1. How you can say that microprocessor integration is helpful to design a client server model?
2. What is distributed computing model, explain with an example?

3. Write about 1 tier and 2 tier model of client server model with examples and their applications in various areas?
4. What is n -tier client server architecture explain with example?
5. How client server model is helpful in division of labor, or how we can distribute various tasks to various machines using client server model?

NOTES

NOTES

CHAPTER

3

TRANSITION TO CLIENT-SERVER PROGRAMMING

STRUCTURE

- 3.1 Transition to client-server programming
- 3.2 Interaction of client and server communication Techniques and protocols
- 3.3 Implementing client server application
- 3.4 Multitasking with process and threads
 - *Summary*
 - *Review Questions*

3.1 TRANSITION TO CLIENT-SERVER PROGRAMMING

One of the greatest challenges facing IT professionals today is how to implement applications using new technology and maintain momentum in the Client Server and Web arenas without negatively impacting the application development staff's ability to stay on top of an overwhelming backlog of work and maintain legacy systems.

A rapidly changing business environment generates a demand for enterprise-wide data access, which, in turn, sets the stage for end user productivity gains. Data access requirements have given rise to an environment in which computers work together to form a system, often called distributed computing, cooperative computing, and the like.

To be competitive in a global economy, organizations in developed economies must employ technology to gain the efficiency necessary to offset their higher labour costs. Re-engineering the business process to provide

NOTES

information and decision-making support at points of customer contact reduces the need for layers of decision-making management, improves responsiveness, and enhance customer service. Empowerment means that knowledge and responsibility are available to the employee at the point of customer contact.

Empowerment will ensure that product and services problems and opportunities are identified and centralized. Client/Server computing is the most effective source for the tools that empower employees with authority and responsibility.

However, Client/Server computing has become more practical and cost-effective because of changes in computer technology that allow the use of PC-based platforms with reliability and robustness comparable to those of traditional mainframe system. In fact, the accelerating trend toward system development based on Internet Technologies, particularly those supplied by Web, has extended the Client/Server model's reach and relevance considerably.

For example, to remain competitive in a global business environment, businesses are increasingly dependent on the Web to conduct their marketing and service operations.

Such Web-based electronic commerce, known as E-commerce, is very likely to become the business norm for businesses of all sizes.

Even a cursory examination of Websites will demonstrate the Web's search. Organizations that range in size from Microsoft, IBM, GM, and Boeing to local arts/craft and flower shops conduct part – or even most – of their business operations via E-commerce. There are various forces that drive the move to client/server computing. Some of them are:

- (i) The changing business environment.
- (ii) Globalization: The world as a market.
- (iii) The growing need for enterprise data access.
- (iv) The demand for end user productivity gains based on the efficient use of data resources.
- (v) Technological advances that have made client/server computing practical like microprocessor technology, data communication and Internet, Database systems, Operating Systems and Graphical User Interface, PC-based and end user application software.
- (vi) Growing cost and performance advantages of PC-based platforms.
- (vii) Enterprise network management.

DRIVING FORCES

Forces that drives the move to Client/Server computing widely can be classified in two general categories based on:

- (i) Business perspective.
- (ii) Technology perspective.

Business Perspective

Basically the business perspective should be kept in mind for obtaining the following achievements through the system:

NOTES

- For increased productivity.
- Superior quality.
- Improved responsiveness.
- Focus on core business.

The effective factors that govern the driving forces are given below:

The changing business environment

Business process engineering has become necessary for competitiveness in the market which is forcing organizations to find new ways to manage their business, despite fewer personnel, more outsourcing, a market driven orientation, and rapid product obsolescence.

Due to globalization of business, the organizations have to meet global competitive pressure by streamlining their operations and by providing an ever-expanding array of customer services. Information management has become a critical issue in this competitive environment; marketing fast, efficient, and widespread data access has become the key to survival. The corporate database has become a far more dynamic asset than it used to be, and it must be available at relatively low cost. Unfortunately, the demand for a more accessible database is not well-served by traditional methods and platforms. The dynamic information driven corporate worlds of today require data to be available to decision makers on time and in an appropriate format. Because end users have become active in handling their own basic data management and data analysis, the movement towards freedom of data access has made Client/Server computing almost inevitable.

One might be tempted to urge that microcomputer networks constitute a sufficient answer to the challenge of dynamic data access. Unfortunately, even the use of networks that tie legions of PC's together is an unsatisfactory solution if request processing overloads the network. The Client/Server model's ability to share resources efficiently by splitting data processing yields a more efficient utilization of those resources. It is not surprising that Client/Server computing has received so much attention from such a wide spectrum of interested parties.

Globalization

Conceptually, the world has begun to be treated as a market. Information Technology plays an important role in bringing all the trade on a single platform by eliminating the barriers. IT helps and supports various marketing priorities like quality, cost, product differentiation and services.

The growing need for enterprise data access

One of the major MIS functions is to provide quick and accurate data access for decision-making at many organizational levels.

NOTES

Managers and decision makers need fast on-demand data access through easy-to-use interfaces. When corporations grow, and especially when they grow by merging with other corporations, it is common to find a mixture of disparate data sources in their systems. For example, data may be located in flat files, in hierarchical or network databases or in relational databases. Given such a multiple source data environment, MIS department managers often find it difficult to provide tools for integrating and aggregating data for decision-making purposes, thus limiting the use of data as a company asset. Client server computing makes it possible to mix and match data as well as hardware. In addition, given the rapidly increasing internet-enabled access to external data through the Internet's inherent Client/Server architecture, corporate Client/Server computing makes it relatively easy to mix external and internal data.

The demand for end user productivity gains based on the efficient use of data resources

The growth of personal computers is a direct result of the productivity gains experienced by end-users at all business levels. End user demand for better ad-hoc data access and data manipulation, better user interface, and better computer integration helped the PC gain corporate acceptance. With sophisticated yet easy to use PCs and application software, end user focus changed from how to access the data to how to manipulate the data to obtain information that leads to competitive advantages.

3.2 INTERACTION OF CLIENT AND SERVER COMMUNICATION TECHNIQUES AND PROTOCOLS

The *client-server* characteristic describes the relationship of cooperating programs in an application. The server component provides a function or service to one or many clients, which initiate requests for such services.

Functions such as email exchange, web access and database access, are built on the client-server model. Users accessing banking services from their computer use a web browser client to send a request to a web server at a bank. That program may in turn forward the request to its own database client program that sends a request to a database server at another bank computer to retrieve the account information. The balance is returned to the bank database client, which in turn serves it back to the web browser client displaying the results to the user. The client-server model has become one of the central ideas of network computing. Many business applications being written today use the client-server model. So do the Internet's main application protocols, such as HTTP, SMTP, Telnet, and DNS.

The interaction between client and server is often described using sequence diagrams. Sequence diagrams are standardized in the Unified

Modeling Language. Specific types of clients include web browsers, email clients, and online chat clients.

Specific types of servers include web servers, ftp servers, application servers, database servers, name servers, mail servers, file servers, print servers, and terminal servers. Most web services are also types of servers.

NOTES

Client-server Interaction

Characteristics of Clients and Servers

In general, client software has the following characteristics.

- It is an application program that becomes a client temporarily when remote access is needed, but performs other computation locally.
- It is invoked by a user and executes for one session.
- It runs locally on the user's computer.
- It actively initiates contact with a server (*CONNECT* primitive).
- It can access multiple services as needed.

In general, server software has the following characteristics.

- It is a special-purpose program dedicated to providing one service.
- It is invoked automatically when a system boots, and continues to execute through many sessions.
- It runs on a shared computer.
- It waits passively for contact from arbitrary remote clients (*LISTEN* primitive).
- It accepts contact from arbitrary clients, but offers a single service.

Note that the word *server* is (strictly) referring to a piece of software. However, a computer running one or more servers is often (incorrectly) called a *server*.

Like most application programs, a client and a server use a transport protocol to communicate. Figure 3.1 illustrates a client and a server using the TCP/IP protocol stack.




Fig. 3.1 : A client and a server

NOTES

The client and server each interact with a protocol in the transport layer of the stack. A sufficiently powerful computer can run multiple servers and clients at the same time. Such a computer must have the necessary hardware resources (e.g. a fast CPU and sufficient memory) and have an operating system capable of running multiple applications concurrently (e.g. UNIX or Windows). Figure 3.2 illustrates such a setup.




Fig. 3.2 : Two servers

The computer in the middle might be running an FTP server and a WWW server. Modern computers are often capable of running many servers at the same time.

Examples of a Client and a Server

Figure 3.3 illustrates the sequence of socket procedure calls required for correct client-server programming.

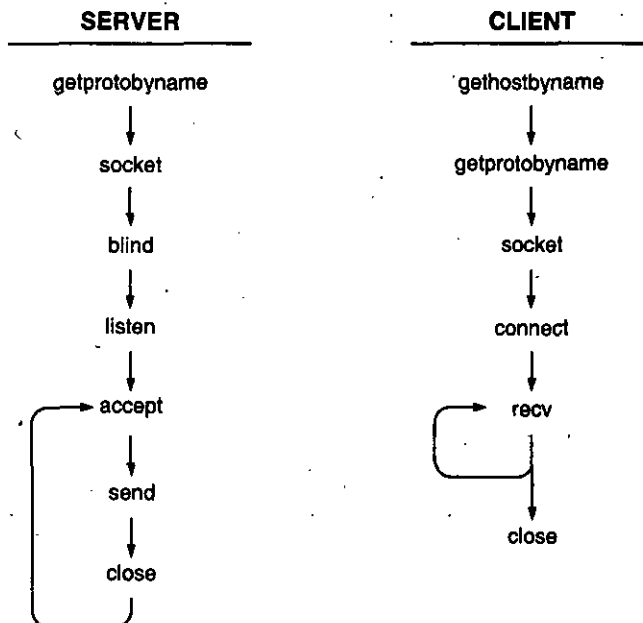


Fig. 3.3 : Socket procedure calls

NOTES

Client and server applications can use either connection-oriented or connectionless transport protocols to communicate. High-level languages, such as Perl can hide this low-level code to make network programming less error-prone.

Example Server

Figure 3.4 illustrates a TCP server application written in Perl. The single function call `IO::Socket::INET->new(...)` executes the first four procedure calls shown in Fig. 3.4

```

use IO::Socket::INET;
use strict;
my $port = shift
    or die "Missing port number\n";
my $socket = IO::Socket::INET->new('LocalPort'=> $port,
                                   'Proto' => 'tcp',
                                   'Listen' => SOMAXCONN)
    or die "Can't create socket ($!)\n";
print "Server listening\n";
while (my $client = $socket->accept) {
    my $name = gethostbyaddr($client->peeraddr, AF_INET);
    my $port = $client->peerport;
    while (<$client>) {
        print "[$name $port] $_";
        print $client "$.: $_";
    }
    close $client
        or die "Can't close ($!)\n";
}
die "Can't accept socket ($!)\n";

```

Fig. 3.4: Simple TCP server

The first two lines specify packages to be used. The next two lines extract a port number from the command line. The next four lines perform a passive open with the server listening for incoming connections. The outer *while* loop accepts a connection and computes the host name and port number of the client. The inner *while* loop reads from the client (`<$client>`), echoes the data on standard output (`print "[$name $port] $_"`), and then sends the client a copy of its own data, prepended with the current line number (`print $client "$.: $_"`). When the client sends end-of-file, the server closes the current connection and continues the outer loop, waiting for another incoming connection.

Client Example

Figure 3.5 illustrated a TCP (connection-oriented) client application written in Perl. The single function call `IO::Socket::INET->new(...)` executes the first four procedure calls shown in Fig. 3.5.

NOTES

```
use IO::Socket::INET;
use strict;
my $name = shift
    or die "Missing server name\n";
my $port = shift
    or die "Missing port number\n";
my $socket = IO::Socket::INET->new('PeerAddr' => $name,
                                   'PeerPort' => $port,
                                   'Proto' => 'tcp')
    or die "Can't create socket ($!)\n";
print "Client sending\n";
while (<STDIN>) {
    print $socket $_;
    print scalar <$socket>;
}
close $socket
    or die "Can't close socket ($!)\n";
```

Fig. 3.5 : Simple TCP client

The first two lines specify packages to be used. The next four lines extract a server name and port number from the command line. The next four lines perform an active open on the socket between our client and the server. The while loop reads standard input (`<STDIN>`), sends the data to the server (`print $socket $_`), and then reads the server's reply (`print scalar <$socket>`). At end-of-file, the connection is closed.

Running the Server

Figure 3.6 illustrates the TCP server application running on penguin and listening on port 1234.

```
penguin(100)% perl tcp-server.pl 1234
Server listening
[localhost 35577] Here we go...
[localhost 35577] Wow, a response!
[localhost 35577] Goodbye

penguin(101)%
```

Fig. 3.6 : Running server

It echoes the segments from a client running on localhost (i.e. the local machine) using port number 35577. It was then terminated.

Running the Client

NOTES

Figure 3.7 illustrates the TCP client application contacting a server running on penguin and listening on port 1234.

```
penguin(107)% perl tcp-client.pl localhost 1234
Client sending
Here we go...
1: Here we go...
Wow, a response!
2: Wow, a response!
Goodbye
3: Goodbye
penguin(108)%
```

Fig. 3.7: Running client

3.3 IMPLEMENTING CLIENT SERVER APPLICATION

Client-server describes an application architecture in which the client requests an action or service from the provider of service, the server. Consider a Web browser and a Web server. When you address a URL in the browser window, it (client) requests a page from a Web server. The server returns an html page to the client, which parses the page (data) and displays it on your computer.

When developing a client-server application, like the Web browser and Web server, you need to consider how you are going to handle developing your application in a team environment and how you are going to handle long-term maintenance. Developing client-server applications parallels developing modular programs. Modular programming separates large applications into smaller constituent pieces to ease development in teams and provide better maintainability. In a client-server application, a module does not have to be part of the same program or even run on the same computer. Each modular function can run on a different device.

How does a client perform

Client programs request service from a server by sending it a message. Referring back to the Web example, a Web browser is a client we use everyday to request Web pages. For example, when you clicked the link to read this article, your browser sent a message to a Web server in Austin, TX. In response, your browser received the html page you are now reading. A Web browser represents many client programs, which manage the graphical user interface (GUI) or display portion of an application; determining the presentation of the service provided by an application.

What is a server's function?

Server programs process client requests by performing the tasks requested by clients. For example, in a Web browser the Web server returns the html page requested by the client. But client requests and server programs are not always so simple. Consider a more complicated application in which you buy a product on a Web page. In this case, the client informs the server what you are purchasing and the server updates a database with the purchase request. Then, the server informs the client that the order has been placed.

Servers are generally passive as they wait for a client request. During these waiting periods servers can perform other tasks or perform maintenance. Unlike the client, the server must continually run because clients can request service at any time. Clients on the other hand only need to run when they require service. Many server applications allow for multiple clients to request service. For example, while you are reading this page others interested in client-server programming could also request and read the same Web page.

Designing a Temperature Monitoring Application

With a basic understanding of client-server applications, how do you implement this technology for your needs? We touch on a simple application of acquiring temperature readings from a thermocouple. First we break this application into three modules making the project easier to develop and maintain. The first module acquires the readings from the thermocouple, the second module analyzes the raw data and relates it to temperature, and the third module provides a graphical user interface or display. Each module must provide data in a format that is understood by the next module.

After dividing the project into manageable pieces, you have a couple of options. You can write one program to perform all three of these functions on the same computer, or you can distribute the process over multiple machines by running each of the individual routines or functions on different machines. At this point, you need to consider your goals. For example, you may need a client-server architecture if you must access the temperature remotely from a different computer than is collecting the temperature data. In this situation, you can create an architecture in which one computer acquires and analyzes the data and another computer or client queries the data to display it. Fig. 3.8 illustrates how the client and server interact.

NOTES

NOTES

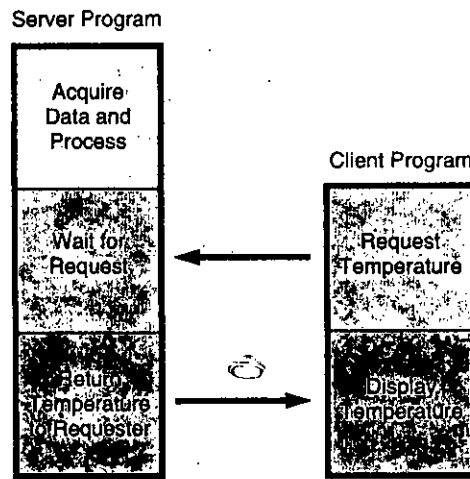


Fig. 3.8

The actual architecture of the application depends on your particular needs. If you are acquiring a large number of data points, it may be better to run the acquisition and analysis modules on separate computers. Thus the display client would request the temperature from the computer processing the raw data, reducing interruptions to the data acquisition process.

Server Considerations

In a temperature monitoring application there are certain considerations for designing the server. These include:

- Is there more than one client?
- Are multiple clients served one-at-a-time or simultaneously?
- Are different clients prioritized or treated differently?

In the simplest application, only one remote client requests the temperature from the temperature acquisition computer. In this case the server waits for a request and returns the temperature to the requesting client.

If multiple clients are serviced, the server continually monitors for service requests. The server queues each request as it is received, ensuring all the clients receive the temperature data they request. The application queues each client and services it in a first-come-first-serve (First In First Out—FIFO) manner. Once the queue is empty, the server continues to wait for new clients. The monitoring process is best broken into two components. One component detects the requests and queues up each client. The second component processes each queued client and performs the task requested one at a time. In this case, the server returns the temperature.

Processing clients one-at-a-time may be suitable for a simple temperature monitoring application. Simultaneously servicing clients is important for more complicated requests that can take a longer time, or if you expect a large number of clients. However, client requests still need to be queued however you design the component that processes the actual request. Servicing

each client in different and independent threads or processes can accomplish this.

Some applications require that you give different clients different priorities. A temperature monitoring system can have passive clients, which display the temperature at any given time, and active clients that query the temperature and process it as part of a larger application. You may need to prioritize the active clients over the passive clients. To do this the server must have a method to identify, differentiate, and authenticate the clients requesting service.

NOTES

To be serviced effectively, the client program must work with the interface or protocol defined by the server. The client program for each of the application architectures described is very similar. In each case the client program must send a temperature request message and must understand the reply from the server. In the case of a server servicing a single remote client or a server servicing multiple clients in a FIFO manner, the client programs have no noticeable difference in architecture. However, a client program interacting with a server that prioritizes different clients has a request mechanism that includes parameters for identification and authentication. This could be as simple as adding two parameters to the temperature request message for an ID and a password that the server can verify.

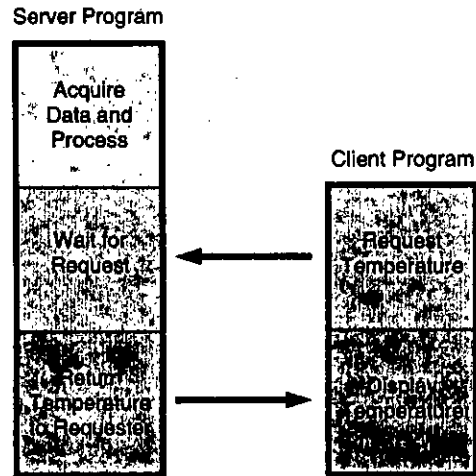
The server defines the interface for communication between the client and server programs. A well designed interface only exposes how to send messages, what messages are supported, and what the response to each message is to the client. The server is also developed in a modular fashion, separating the message handling mechanism from the functional component processing the request itself. Therefore, you can update the functional component of the server without affecting any clients. This modular approach to a client-server application provides a distributed application for remote temperature monitoring that you can maintain and update without affecting each client and server component.

In Part I, we discussed the fundamentals of client-server applications through a simple temperature monitoring application. In Part II, we use LabVIEW code samples to review these principles and learn how to develop client-server applications.

Client-server applications must perform the primary functions discussed in previous section: client programs request service from a server by sending it a message, and server programs process client requests by performing the tasks requested by clients.

With this basic understanding of client-server applications, let us look more closely at a temperature monitoring application. The server in this application continually detects requests from client applications, processes the requests, and responds by returning the current temperature to the client. Refer to the end of this article for the LabVIEW code illustrated in this article.

NOTES



Please note that using TCP/IP methods may not be the most effective way to develop applications where data is being shared or distributed. You can develop applications faster and with built-in security options using simpler methods and features in LabVIEW, including DataSocket and VI Server. This article uses TCP/IP methods for the purpose of illustrating the concepts of client-server applications.

A Simple LabVIEW Client-Server Application

In previous section we defined a simple temperature monitoring application. We divided the application into three modules to make the application easier to maintain. Using while loops, you can implement the major components of the server application. These components are—detecting temperature data requests, processing temperature data requests, and responding to requests. Your first loop executes until the server receives the connection request. The second loop continues to provide temperature data to the client until the client disconnects from the server.

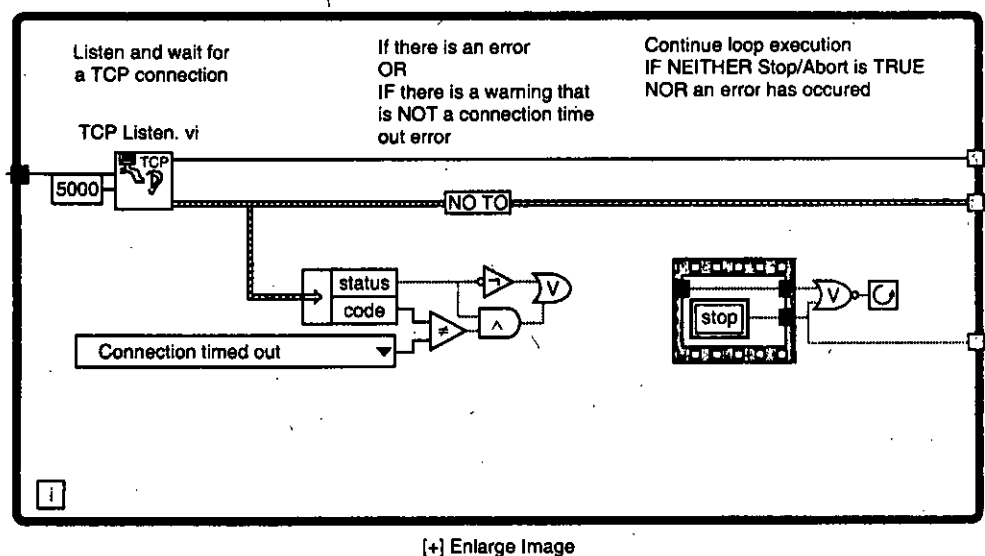


Fig. 3.9 : Server Waiting for a Client Request

Before we give details, let us plan the client-server steps in our simple temperature monitoring application. The client and server follow these steps:

1. Client sends a request.
2. Server detects the request using the TCP Listen.
3. Client connects to the server using the TCP Open Connection.
4. Server uses the TCP Write.vi to send temperature data to the client.
5. Client reads the transmitted data using the TCP Read.

More specifically, here are the actions the server takes when detecting a client request. The server detects a client request with the TCP Listen.vi, found in the TCP utility library (<labview-dir>\vi.lib\Utility\tcp.llb). The VI creates a listener which waits for an accepted TCP connection from a client. This is similar to the way your telephone operates. The telephone in your house waits or listens for a call or connection from another phone. The TCP Listen.vi enables your server to wait or listen for a connection from a server. The loop, illustrated in Fig. 3.10, waits for a connection from a client and processes any errors or warnings from the TCP Listen.vi. Error checking and support of an abort mechanism is not necessary for the application, however it is generally good programming practice because you can handle different errors appropriately and the server operator can quit the application in an emergency.

Once the server detects the client request, the client connects to the server with the TCP Open Connection.vi, found in the *Communication-TCP* function palette. The client connects to a specific server application by specifying the server address—the IP address or hostname—and the port.




Fig. 3.10 : Client Opening a Connection

When the connection between the client and server is made, the server sends temperature data to the client, and the client reads and displays the returned data. TCP Write.vi and TCP Read.vi primarily pass data between the client and server.

In Figs. 3.11 and 3.12 below, we illustrate the associated code for the data transfer between the client-server applications.

NOTES

NOTES

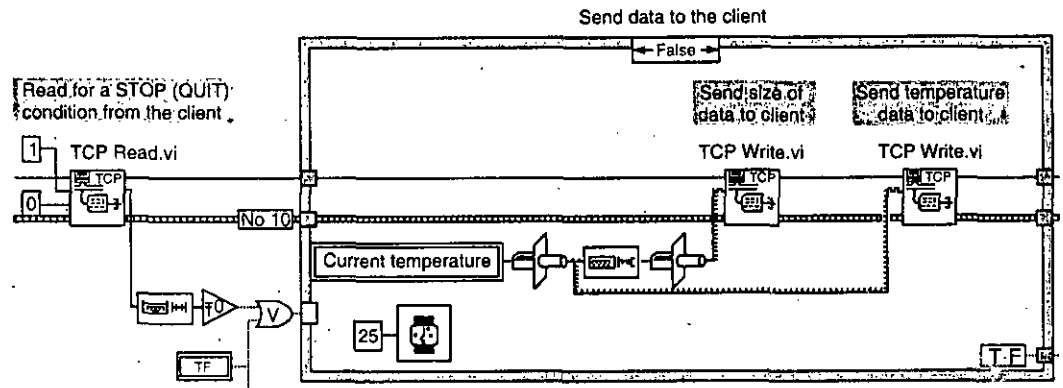


Fig. 3.11 : Server Sending Temperature Data to the Client




Fig. 3.12 : Client Reading Temperature Data from the Server

Our temperature monitoring application is fairly simple. However, in larger applications, the size of the data you request from the server can put limitations on your application. Therefore, before sending the actual temperature data, the server sends the size of temperature data being transferred. This mechanism is necessary in applications where a server sends an array or other complex data structure of unknown size. This also ensures that no data is missed.

Client-Server Communication

Client-server applications are not limited to transferring data. Many times, you need to run an application that synchronizes the client-server application, indicates the state, or identifies and authenticates clients. If you need this type of data for your application, use the TCP Read.vi and TCP Write.vi to read and write data between the client and server. You must predefine any information you need to provide this added functionality.

Streaming Temperature Data

In a simple situation, you may require a client to query for each temperature point one at a time. In Figs. 3.13 and 3.14, the client connects

once to the server and continuously receives temperature data until the client is stopped.

Transition to Client-Server Programming

The client stops the data transfer by sending a byte of data ("Q") to the server (Fig. 3.14). The server tries to read data the client sends. If it detects that only one byte of information has been sent, the server knows the client has requested a termination of the connection. If you did not predetermine this relationship, the server would continue to send temperature data to the client.

NOTES

Authenticating/Validating Clients

Some applications and environments require that you authenticate a client before a server can return data. If you are concerned about sensitive or confidential information, it is important to only service authorized clients. You can validate clients as a separate module to easily incorporate it into the client-server relationship.




Fig. 3.13 : Validating Client Access to Server Data

You can authenticate clients using an access control list, or a list of authorized clients, on the server machine. The TCP/IP Access Validation—refer to the TCP/IP Access Validation link below—checks clients against this list to validate them.

Alternatively, you can use passwords to validate a client. After connecting to the server, the client sends a password and other identifiable information to the server using TCP Write.vi. The validation module in the server reads the password sent by the client and validates it. If the password is incorrect, the server denies access to the client.

NOTES




Fig. 3.14 : Password Validation Module

In Fig. 3.14, we illustrate a simple password authentication mechanism. Here, the client provides a password which is verified against a list of acceptable passwords maintained by the server.

Security Considerations

We have covered basic client validation. However, what if you have a more in-depth security concern? How do you implement additional security measures in a client-server application? We point out some additional considerations; this is not an exhaustive discussion of security in client-server applications.

When you use access control list or passwords to authenticate a client, you transfer the password in the "open". Refer to the example code at the end of this article. You have not encrypted or protected the password, leaving it open to a malicious application, which can monitor network traffic to obtain the password. The malicious application has compromised the password and can use it to access data from the server. For this reason, it is important to protect passwords and other data used for authentication through encryption, reducing the risk of the password being compromised.

In addition to encrypting passwords, sometimes you need to encrypt all data transferred between the client and server. For example, you may want to encrypt your data if your application is distributed to a client or server outside your company firewall.

Serving Multiple Clients with LabVIEW

Many applications require more than one client to access data from the server. The server must handle multiple connections and respond to each and every client requesting service. The server can respond in a First-In-First Out (FIFO) manner or handle multiple clients simultaneously—in parallel.

When you try to handle multiple clients with one server, you can realize the benefits of modular programming. Your overall structure remains similar to those we have already discussed. Let's refer back to our temperature monitoring application. When serving multiple clients, you need to separate

Transition to Client-Server Programming

NOTES




Fig. 3.16 uses a while loop to send data to clients and the queue to determine if any client have made requests.

NOTES

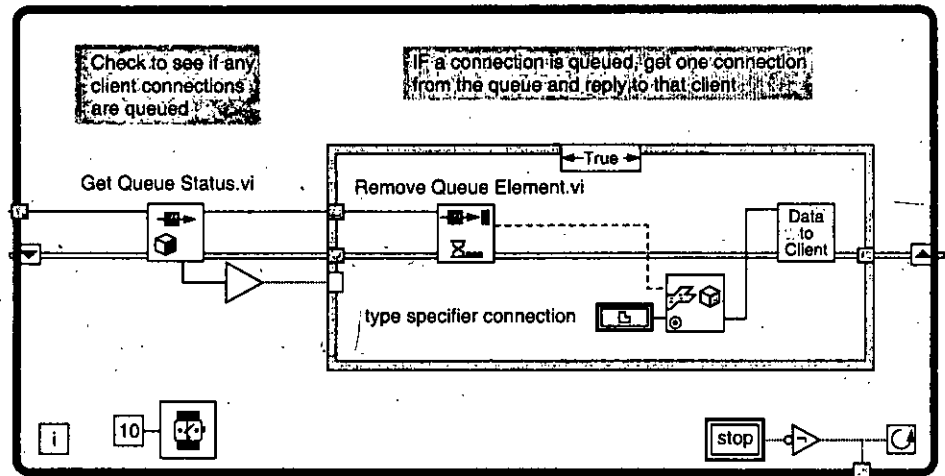


Fig. 3.16 : Handling Queued Client Requests—A While Loop

With this mechanism, the server responds to clients in a FIFO manner. A FIFO response mechanism fails if an individual client request takes a long time, starving other clients of service. In this case, the server must impose a limit on the time an individual client can receive service or handle clients simultaneously or in parallel.

Refer to the Understanding Client-Server Applications—Example Code to investigate a FIFO Server. Refer to "SERVER - FIFO.vi".

Serving Multiple Clients in Parallel

Handling multiple clients in parallel can be tricky in a dataflow environment. The only architectural difference is choosing to serve clients in FIFO or in parallel.

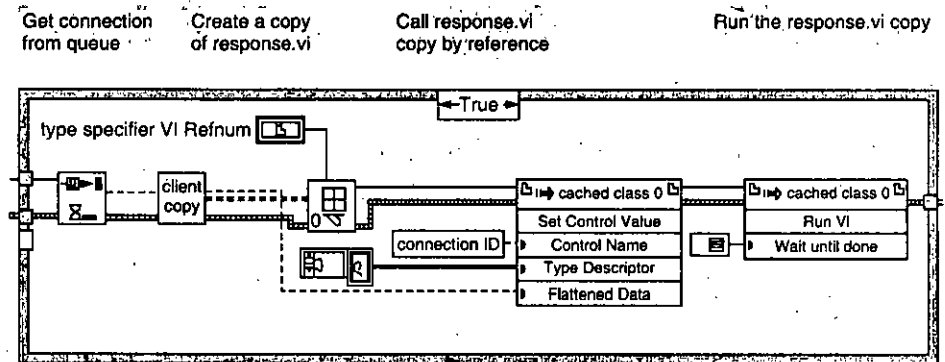


Fig. 3.17 : Handled Queued Client Requests in Parallel

In parallel, the server creates a copy of the SubVI used to respond to a client. This is the same SubVI used in the FIFO server. It then calls this copy by reference and runs it. The copy of the VI responds to the corresponding client and continues to run until the client stops running. The server creates a new copy of the SubVI for each client request. In this

way, it can call each copy independently and handle multiple client requests in parallel.

3.4 MULTITASKING WITH PROCESS AND THREADS

NOTES

What Is Multitasking?

Multitasking is the ability of a computer to run more than one program, or *task*, at the same time. Multitasking contrasts with single-tasking, where one process must entirely finish before another can begin. MS-DOS is primarily a single-tasking environment, while Windows 3.1 and Windows NT are both multi-tasking environments.

On a single-processor multitasking system, multiple processes don't actually run at the same time since there's only one processor. Instead, the processor switches among the processes that are active at any given time. Because computers are so fast compared with people, however, it appears to the user as though the computer is executing all of the tasks at once. Multitasking also allows the computer to make good use of the time it would otherwise spend waiting for I/O devices and user input—that time can be used for some other task that doesn't need I/O at the moment.

Multitasking on a multiple-processor system still involves the processors switching between tasks because there are almost always more tasks to run than there are processors. Note, however, that there can be as many tasks running simultaneously as there are processors in the system. For the moment, we'll discuss multitasking on a single-processor system.

Preemptive and Non-Preemptive Multitasking

Within the category of multitasking, there are two major sub-categories: *preemptive* and *non-preemptive* (or cooperative). In *non-preemptive multitasking*, use of the processor is never taken from a task; rather, a task must voluntarily yield control of the processor before any other task can run. Windows 3.1 uses non-preemptive multitasking for Windows applications.

Programs running under a non-preemptive operating system must be specially written to cooperate in multitasking by yielding control of the processor at frequent intervals. Programs that do not yield sufficiently often cause non-preemptive systems to stay "locked" in that program until it does yield. An example of failed non-preemptive multitasking is the inability to do anything else while printing a document in Microsoft Word for Windows 2.0a. This happens because Word does not give up control of the processor often enough while printing your document. The worst case of a program not yielding is when a program crashes. Sometimes, programs which crash in Windows 3.1 will crash the whole system simply because no other tasks can run until the crashed program yields.

NOTES

Preemptive multitasking differs from non-preemptive multitasking in that the operating system can take control of the processor without the task's cooperation. (A task can also give it up voluntarily, as in non-preemptive multitasking.) The process of a task having control taken from it is called **preemption**. Windows NT uses preemptive multitasking for all processes except 16-bit Windows 3.1 programs. As a result, a Windows NT application cannot take over the processor in the same way that a Windows 3.1 application can.

A preemptive operating system takes control of the processor from a task in two ways:

- When a task's time quantum (or time slice) runs out. Any given task is only given control for a set amount of time before the operating system interrupts it and schedules another task to run.
- When a task that has higher priority becomes ready to run. The currently running task loses control of the processor when a task with higher priority is ready to run regardless of whether it has time left in its quantum or not.

Switching Among Tasks

At any given time, a processor (CPU) is executing in a specific context. This context is made up of the contents of its registers and the memory (including stack, data, and code) that it is addressing. When the processor needs to switch to a different task, it must save its current context (so it can later restore the context and continue execution where it left off) and switch to the context of the new task. This process is called **context switching**.

When Windows NT switches tasks, it saves all of the registers of the task it's leaving and re-loads the registers of the task to which it's switching. This process also enables the proper address space for the task to which Windows NT is switching.

Processes and Threads

In addition to being a preemptive multitasking operating system, Windows NT is also multithreaded, meaning that more than one thread of execution (or **thread**) can execute in a single task at once.

A process comprises:

- A private memory address space in which the process's code and data are stored.
- An access token against which Windows NT makes security checks.
- System resources such as files and windows (represented as object handles).
- At least one thread to execute the code.

A thread comprises:

- A processor state including the current instruction pointer.
- A stack for use when running in user mode.

- A stack for use when running in kernel mode.

Since processes (not threads) own the access token, system resource handles, and address space, threads do not have their own address spaces nor do they have their own access token or system resource handles. Therefore, all of the threads in a process SHARE the same memory, access token, and system resources (including quota limits) on a “per-process” rather than a “per-thread” basis. In a multithreaded program, the programmer is responsible for making sure that the different threads don’t interfere with each other by using these shared resources in a way that conflicts with another thread’s use of the same resource. (As you might suspect, this can get a little tricky.)

Why Use Multithreading?

Multithreading provides a way to have more than one thread executing in the same process while allowing every thread access to the same memory address space. This allows very fast communication among threads. Threads are also easier to create than processes since they don’t require a separate address space.

Inside Windows NT, processes and threads are represented as objects that are created, maintained, and destroyed by the Process Manager. These Process Manager process and thread objects contain simpler kernel process and thread objects.

Some typical examples of the use of multiple threads are using a background thread to print a document in a word processor and to recalculate a spreadsheet. When a new thread is created to do these tasks, the main thread can continue responding to user input. A single-threaded application can’t respond to user input until it’s done printing or recalculating or whatever.

On a uniprocessor platform, the use of multiple threads allows a user to continue using a program even while another thread is doing some lengthy procedure. But only one thread executes at a time.

On a multiprocessor platform, more than one processor may be running different threads in the same process. This has the potential for very significantly speeding up the execution of your program.

Sharing A Single Address Space—Synchronizing Access To Data

Running each process in its own address space had the advantage of reliability since no process can modify another process’s memory. However, all of a process’s threads run in the same address space and have unrestricted access to all of the same resources, including memory. While this makes it easy to share data among threads, it also makes it easy for threads to step on each other. As mentioned before, multithreaded programs must be specially programmed to ensure that threads don’t step on each other.

A section of code that modifies data structures shared by multiple threads is called a **critical section**. It is important that when a critical

NOTES

NOTES

section is running in one thread that no other thread be able to access that data structure. Synchronization is necessary to ensure that only one thread can execute in a critical section at a time. This synchronization is accomplished through the use of some type of Windows NT synchronization object. Programs use Windows NT synchronization objects rather than writing their own synchronization both to save coding effort and for efficiency, when you wait on a Windows NT synchronization object, you do not use any CPU time testing the object to see when it's ready.

Windows NT provides a variety of different types of synchronization objects that programs can use to coordinate threads' access to shared data structures. Synchronization objects remember their states and can be set and tested in one uninterruptable step. They also cause the thread to be suspended while waiting on an object and to automatically restart when the other thread signals that it's done.

During the initialization of a program, the program creates a synchronization object for each data structure or object that will be shared among threads.

Every critical section will have the following structure:

1. Wait on the synchronization object before accessing the data structure. The Windows NT waiting API insures that your thread is suspended until the synchronization object becomes unlocked. As soon as the synchronization object becomes unlocked, Windows NT sets the synchronization object to "locked" and restarts your thread.
2. Access the data structure. (This is the critical section.)
3. Unlock the synchronization object so that the data can be accessed by other threads.

The first step is critical because if it's omitted then any thread can access the data structure while you're accessing. The last step is also critical—it it's omitted, then no thread will be able to access the data even after you're done.

Using this technique on every critical section insures that only one thread can access the data at a time.

The Life Cycle of a Thread

Each thread has a *dispatcher state* that changes throughout its lifetime.

The most important dispatcher states are:

- Running: only one thread per processor can be running at any time.
- Ready: threads that are in the Ready state may be scheduled for execution the next time the kernel dispatches a thread. Which Ready thread executes is determined by their priorities.
- Waiting: threads that are waiting for some event to occur before they become Ready are said to be waiting. Examples of events

include waiting for I/O, waiting for a message, and waiting for a synchronization object to become unlocked.

The Kernel's Dispatcher

The kernel's dispatcher performs scheduling and context switching.

Thread scheduling is the act of determining which thread runs on each processor at a given time.

Context switching is the act of saving one thread's volatile state (CPU register contents) and restoring another thread's state so it can continue running where it previously left off.

How Thread Priorities Affect Scheduling

The kernel's dispatcher schedules threads to run based a 32-level priority scheme. Windows NT guarantees that the threads that are ready that have the highest priority will be running at any given time. (That's one thread on a single-processor system.) Threads with a priority of 31 will be run before any others, while threads with a priority of 0 will run only if no other threads are ready. The range of priorities is divided in half with the upper 16 reserved for real-time threads and the lower 16 reserved for variable priority threads.

Real-time threads run at the same priority for their entire lifetime. They are commonly used to monitor or control systems that require action to be taken at very precise intervals. These threads run at higher priorities than all variable priority threads, which means that they must be used sparingly.

Variable priority threads are assigned a base priority when they are created. (A thread's base priority is determined by the process to which the thread belongs). The priority of such threads can be adjusted dynamically by the kernel's dispatcher. A thread's dynamic priority can vary up to two priority levels above or below its base priority.

The dispatcher maintains a **priority queue** of ready tasks. When prompted to reschedule, it changes the state of the highest priority task to Standby. When the conditions are right, a context switch is performed to begin the thread's execution and the thread goes into the Ready state.

Lower priority threads will always be preempted when a higher priority thread enters the ready state. This is true even if the lower priority thread has time remaining in its quantum, or if the lower priority thread is running on a different processor.

Performance Tuning

In order to get the computer system to perform as users expect, Windows NT changes the priorities of threads over time.

Each process has a base priority. Threads in a process can alter their base priority by up to two levels up or down.

NOTES

Depending on the type of work the thread is doing, Windows NT may also adjust the thread's dynamic priority upwards from its base priority. For instance:

NOTES

- Threads that are waiting for input get a priority boost, as do threads in the foreground process. This makes the system responsive to the user.
- Threads get a priority boost after completing a voluntary wait.
- All threads periodically get a priority boost to prevent lower priority threads from holding locks on shared resources that are needed by higher priority threads.
- Compute-bound threads get their priorities lowered.

Scheduling on Multiprocessor Systems

A **multiprocessing** operating system is one that can run on computer systems that contain more than one processor. Windows NT is a **symmetric multiprocessing (SMP)** system, meaning that it assumes that all of the processors are equal and that they all have access to the same physical memory. Therefore, Windows NT can run any thread on any available processor regardless of what process, user or Executive, owns the thread.

There are also asymmetric multiprocessing (ASMP) systems in which processors are different from each other—they may address different physical memory spaces, or they may have other differences. These operating systems only run certain processes on certain processors—for instance, the kernel might always execute on a particular processor.

The design of Windows NT supports **processor affinity**, whereby a process or thread can specify that it is to run on a particular set of processors, but this facility isn't supported in the first release.

Windows NT uses the same rules for scheduling on a multiprocessor system as it does on a single processor system, so at any given time the threads that are ready and have the highest priorities are actually running.

SUMMARY

1. Client/Server computing has become more practical and cost-effective because of changes in computer technology that allow the use of PC-based platforms with reliability and robustness comparable to those of traditional mainframe system.
2. The client-server characteristic describes the relationship of cooperating programs in an application. The server component provides a function or service to one or many clients, which initiate requests for such services.
3. Client-server describes an application architecture in which the client requests an action or service from the provider of service, the server.
4. Server programs process client requests by performing the tasks requested by clients.
5. Multitasking is the ability of a computer to run more than one program, or task, at the same time.

6. Real-time threads run at the same priority for their entire lifetime.
7. Variable priority threads are assigned a base priority when they are created.

REVIEW QUESTIONS

NOTES

1. What is the difference between preemptive and non-preemptive multitasking? In preemptive multitasking, what is a time quantum?
2. What are processes and threads in Windows NT? Are objects owned by processes or threads? How can an application benefit from having multiple threads of execution?
3. Why will a multithreaded Windows NT application execute faster on a multiprocessor workstation than on a uniprocessor system, while a single-threaded application executes at roughly the same speed on both systems?
4. What is meant by the terms thread scheduling and context switching? What type of information makes up a thread's context?
5. What are the relative priorities of real-time and variable-priority threads? What are each used for? What can cause the priority of a variable-priority thread to change?
6. What event(s) cause a thread to make the following dispatcher state transitions: Running to Waiting, Running to Ready, Waiting to Ready, Ready to Running?
7. If a lower priority thread is running, will it be preempted when a higher priority thread enters the ready state, even if the lower priority thread has time remaining in its quantum?
8. Why is synchronization of threads sometimes necessary? How does this relate to critical sections?
9. What is multiprocessing? What is the difference between symmetric multiprocessing (SMP) and asymmetric multiprocessing (ASMP)? What is processor affinity? How are multiple processors scheduled under Windows NT?
10. Design a client server model for on line shopping mart.

NOTES

CHAPTER

4

CLIENT SERVER SCHEDULING AND MEMORY MANAGEMENT

STRUCTURE

- 4.1 Scheduling implementations
- 4.2 Scheduler internals
- 4.3 Preemptive Vs non-preemptive systems
- 4.4 Synchronization
- 4.5 Semaphores
- 4.6 Semaphore implementation
- 4.7 Memory-management allocation, sharing and manipulating
 - *Summary*
 - *Review Questions*

4.1 SCHEDULING IMPLEMENTATIONS

Scheduling is a key concept in computer multitasking, multiprocessing operating system and real-time operating system designs. **Scheduling** refers to the way processes are assigned to run on the available CPUs, since there are typically many more processes running than there are available CPUs. This assignment is carried out by software known as a **scheduler** and **dispatcher**.

The **scheduler** is concerned mainly with:

- CPU utilization-to keep the CPU as busy as possible.
- Throughput-number of processes that complete their execution per time unit.

- Turnaround-total time between submission of a process and its completion.
- Waiting time-amount of time a process has been waiting in the ready queue.
- Response time-amount of time it takes from when a request was submitted until the first response is produced.
- Fairness-Equal CPU time to each thread.

In real-time environments, such as mobile devices for automatic control in industry (for example robotics), the scheduler also must ensure that processes can meet deadlines; this is crucial for keeping the system stable. Scheduled tasks are sent to mobile devices and managed through an administrative back end.

NOTES

4.2 SCHEDULER INTERNALS

Scheduling disciplines are algorithms used for distributing resources among parties which simultaneously and asynchronously request them. Scheduling disciplines are used in routers (to handle packet traffic) as well as in operating systems (to share CPU time among both threads and processes), disk drives (I/O scheduling), printers (print spooler), most embedded systems, etc.

The main purposes of scheduling algorithms are to minimize resource starvation and to ensure fairness amongst the parties utilizing the resources. Scheduling deals with the problem of deciding which of the outstanding requests is to be allocated resources. There are many different scheduling algorithms. In this section, we introduce several of them.

First in first out

Also known as *First Come, First-Served* (FCFS), its the simplest scheduling algorithm, simply queues processes in the order that they arrive in the ready queue.

- Since context switches only occur upon process termination, and no reorganization of the process queue is required, scheduling overhead is minimal.
- Throughput can be low, since long processes can hog the CPU.
- Turnaround time, waiting time and response time can be low for the same reasons above.
- No prioritization occurs, thus this system has trouble meeting process deadlines.
- The lack of prioritization does permit every process to eventually complete, hence no starvation.

NOTES

Shortest remaining time

Also known as *Shortest Job First* (SJF). With this strategy the scheduler arranges processes with the least estimated processing time remaining to be next in the queue. This requires advanced knowledge or estimations about the time required for a process to complete.

- If a shorter process arrives during another process' execution, the currently running process may be interrupted, dividing that process into two separate computing blocks. This creates excess overhead through additional context switching. The scheduler must also place each incoming process into a specific place in the queue, creating additional overhead.
- This algorithm is designed for maximum throughput in most scenarios.
- Waiting time and response time increase as the process' computational requirements increase. Since turnaround time is based on waiting time plus processing time, longer processes are significantly affected by this. Overall waiting time is smaller than FIFO, however since no process has to wait for the termination of the longest process.
- No particular attention is given to deadlines, the programmer can only attempt to make processes with deadlines as short as possible.
- Starvation is possible, especially in a busy system with many small processes being run.

Fixed priority pre-emptive scheduling

The O/S assigns a fixed priority rank to every process, and the scheduler arranges the processes in the ready queue in order of their priority. Lower priority processes get interrupted by incoming higher priority processes.

- Overhead is not minimal, nor is it significant.
- FPPS has no particular advantage in terms of throughput over FIFO scheduling.
- Waiting time and response time depend on the priority of the process. Higher priority processes have smaller waiting and response times.
- Deadlines can be met by giving processes with deadlines a higher priority.
- Starvation of lower priority processes is possible with large amounts of high priority processes queuing for CPU time.

Round-robin scheduling

The scheduler assigns a fixed time unit per process, and cycles through them.

- RR scheduling involves extensive overhead, especially with a small time unit.
- Balanced throughput between FCFS and SJN, shorter jobs are completed faster than in FCFS and longer processes are completed faster than in SJN.

- Fastest average response time, waiting time is dependent on number of processes, and not average process length.
- Because of high waiting times, deadlines are rarely met in a pure RR system.
- Starvation can never occur, since no priority is given. Order of time unit allocation is based upon process arrival time, similar to FCFS.

NOTES

Multilevel queue scheduling

This is used for situations in which processes are easily classified into different groups. For example, a common division is made between foreground (interactive) processes and background (batch) processes. These two types of processes have different response-time requirements and so may have different scheduling needs.

Process states

In a multitasking computer system, processes may occupy a variety of states. These distinct states may not actually be recognized as such by the operating system kernel, however they are a useful abstraction for the understanding of processes.

The following typical process states are possible on computer systems of all kinds. In most of these states, processes are "stored" on main memory.

Created

(Also called **New**) When a process is first created, it occupies the "created" or "new" state. In this state, the process awaits admission to the "ready" state. This admission will be approved or delayed by a long-term, or admission, scheduler. Typically in most desktop computer systems, this admission will be approved automatically, however for real-time operating systems this admission may be delayed. In a real time system, admitting too many processes to the "ready" state may lead to oversaturation and over contention for the systems resources, leading to an inability to meet process deadlines.

Ready or Running

(Also called **waiting** or **runnable**) A "ready" or "waiting" process has been loaded into main memory and is awaiting execution on a CPU (to be context switched onto the CPU by the dispatcher, or short-term scheduler). There may be many "ready" processes at any one point of the systems execution-for example, in a one processor system, only one process can be executing at any one time, and all other "concurrently executing" processes will be waiting for execution.

A *ready queue* is used in computer scheduling. Modern computers are capable of running many different programs or processes at the same time. However, the CPU is only capable of handling one process at a time. Processes that are ready for the CPU are kept in a queue for "ready"

processes. Other processes that are waiting for an event to occur, such as loading information from a hard drive or waiting on an internet connection, are not in the ready queue.

NOTES

Blocked

A process that is waiting for some event (such as I/O operation completion or a signal).

Terminated

A process may be terminated, either from the "running" state by completing its execution or by explicitly being killed. In either of these cases, the process moves to the "terminated" state. If a process is not removed from memory after entering this state, this state may also be called **zombie**.

Context Switching

A *context switch* (also sometimes referred to as a *process switch* or a *task switch*) is the switching of the CPU (central processing unit) from one *process* or *thread* to another.

A process (also sometimes referred to as a *task*) is an *executing* (i.e., running) instance of a program. In Linux, threads are lightweight processes that can run in parallel and share an *address space* (i.e., a range of memory locations) and other resources with their *parent processes* (i.e., the processes that created them).

A *context* is the contents of a CPU's *registers* and *program counter* at any point in time. A register is a small amount of very fast memory inside of a CPU (as opposed to the slower RAM main memory outside of the CPU) that is used to speed the execution of computer programs by providing quick access to commonly used values, generally those in the midst of a calculation. A *program counter* is a specialized register that indicates the position of the CPU in its instruction sequence and which holds either the address of the instruction being executed or the address of the next instruction to be executed, depending on the specific system.

Context switching is an essential feature of *multitasking* operating systems. A multitasking operating system is one in which multiple processes execute on a single CPU seemingly simultaneously and without interfering with each other. This illusion of *concurrency* is achieved by means of context switches that are occurring in rapid succession (tens or hundreds of times per second). These context switches occur as a result of processes voluntarily relinquishing their time in the CPU or as a result of the *scheduler* making the switch when a process has used up its *CPU time slice*.

A context switch can also occur as a result of a *hardware interrupt*, which is a signal from a hardware device (such as a keyboard, mouse, modem or system clock) to the kernel that an *event* (e.g., a key press, mouse movement or arrival of data from a network connection) has occurred.

Critical section

In concurrent programming a **critical section** is a piece of code that accesses a shared resource (data structure or device) that must not be concurrently accessed by more than one thread of execution. A critical section will usually terminate in fixed time, and a thread, task or process will have to wait a fixed time to enter it (aka bounded waiting). Some synchronization mechanism is required at the entry and exit of the critical section to ensure exclusive use, for example a semaphore.

By carefully controlling which variables are modified inside and outside the critical section (usually, by accessing important state only from within), concurrent access to that state is prevented. A critical section is typically used when a multithreaded program must update multiple related variables without a separate thread making conflicting changes to that data. In a related situation, a critical section may be used to ensure a shared resource, for example a printer, can only be accessed by one process at a time.

Mutual exclusion

Mutual exclusion (often abbreviated to **mutex**) algorithms are used in concurrent programming to avoid the simultaneous use of a common resource, such as a global variable, by pieces of computer code called critical sections. A critical section is a piece of code in which a process or thread accesses a common resource. The critical section by itself is not a mechanism or algorithm for mutual exclusion. A program, process, or thread can have the critical section in it without any mechanism or algorithm which implements mutual exclusion.

Mutual exclusion is simply defined as program execution on a set of resources for which access must only be granted to one task at a time. Tasks are said to be mutually exclusive if they cannot execute the same operations concurrently.

There are many ways to provide mutual exclusion, including locking variables, disabling interrupts and using semaphores.

4.3 PREEMPTIVE VS NON-PREEMPTIVE SYSTEMS

Preemptive scheduling allows a process to be interrupted in the midst of its execution, taking the CPU away and allocating it to another process.

Non-preemptive scheduling ensures that a process relinquishes control of the CPU only when it finishes with its current CPU burst.

Preemptive systems

The term **preemptive multitasking** is used to distinguish a multitasking operating system, which permits preemption of tasks, from a cooperative multitasking system wherein processes or tasks must be explicitly programmed to yield when they do not need system resources.

NOTES

NOTES

In simple terms: Preemptive multitasking involves the use of an interrupt mechanism which suspends the currently executing process and invokes a scheduler to determine which process should execute next. Therefore all processes will get some amount of CPU time at any given time.

In preemptive multitasking, the operating system kernel can also initiate a context switch to satisfy the scheduling policy's priority constraint, thus preempting the active task. In general, preemption means "prior seizure of". When the high priority task at that instance seizes the currently running task, it is known as preemptive scheduling.

The term "preemptive multitasking" is sometimes mistakenly used when the intended meaning is more specific, referring instead to the class of scheduling policies known as *time-shared scheduling*, or *time-sharing*.

Preemptive multitasking allows the computer system to more reliably guarantee each process a regular "slice" of operating time. It also allows the system to rapidly deal with important external events like incoming data, which might require the immediate attention of one or another process.

At any specific time, processes can be grouped into two categories: those that are waiting for input or output (called "I/O bound"), and those that are fully utilizing the CPU ("CPU bound"). In early systems, processes would often "poll", or "busywait" while waiting for requested input (such as disk, keyboard or network input). During this time, the process was not performing useful work, but still maintained complete control of the CPU. With the advent of interrupts and preemptive multitasking, these I/O bound processes could be "blocked", or put on hold, pending the arrival of the necessary data, allowing other processes to utilize the CPU. As the arrival of the requested data would generate an interrupt, blocked processes could be guaranteed a timely return to execution.

Although multitasking techniques were originally developed to allow multiple users to share a single machine, it soon became apparent that multitasking was useful regardless of the number of users. Many operating systems, from mainframes down to single-user personal computers and no-user control systems (like those in robotic spacecraft), have recognized the usefulness of multitasking support for a variety of reasons. Multitasking makes it possible for a single user to run multiple applications at the same time, or to run "background" processes while retaining control of the computer.

Time slice

The period of time for which a process is allowed to run in a preemptive multitasking system is generally called the *time slice*, or *quantum*. The scheduler is run once every time slice to choose the next process to run. If the time slice is too short then the scheduler will consume too much processing time.

An interrupt is scheduled to allow the operating system kernel to switch between processes when their time slices expire, effectively allowing

the processor's time to be shared between a number of tasks, giving the illusion that it is dealing with these tasks simultaneously, or concurrently. The operating system which controls such a design is called a multi-tasking system.

Systems supporting preemptive multitasking

Today, nearly all operating systems support preemptive multitasking, including the current versions of Windows, Mac OS, Linux, iOS and Android. Examples of early preemptive operating systems include AmigaOS, the Windows 95/98/ME (32-bit applications only) and Windows NT family (including 2000, XP, Vista, and 7), Linux, BSD, OS/2 (2.x onwards), Mac OS X. Unix and Unix-like systems, VMS, and other systems used in the academic and medium-to-large business markets, have always supported preemptive multitasking.

Examples of older, non-preemptive (cooperative) operating systems include Windows 1.x, 2.x, 3.x, Windows for Workgroups, Windows 95/98/ME (when running 16-bit applications), NetWare, and Classic Mac OS versions (System 5 and up). Non-multitasking operating systems include older versions of Mac OS, MS-DOS, and Commodore 64 OS which could only execute one program at a time.

Some of the earliest operating systems available to home users featuring preemptive multitasking were Sinclair QDOS (1984) and Amiga OS (1985). These both ran on Motorola 68000-family microprocessors without memory management. Amiga OS used dynamic loading of relocatable code blocks ("hunks" in Amiga jargon) to multitask preemptively all processes in the same flat address space. Other Unix-like systems including MINIX and Coherent provided preemptive multitasking on 1980s-era personal computers.

Early PC operating systems, such as MS-DOS and DR-DOS, did not support multitasking at all. Novell NetWare, Microsoft Windows and OS/2 systems introduced cooperative multitasking to the PC, but did not support preemptive multitasking. In the case of the PC, the slow start was partly because of the need to support a large legacy code of DOS software written to run in single-user mode, whereas the Amiga system was designed to multitask from the beginning. The earliest version of Windows to support a limited form of preemptive multitasking was Windows 2.1x, which used the Intel 80386's Virtual 8086 mode to run DOS applications in virtual 8086 machines—commonly known as "DOS boxes"—which could be preempted. In Windows 95, 98, and ME, 32-bit applications were made preemptive by running each one in a separate address space, but 16 bit applications remained cooperative. ^[2] Windows NT has always supported preemptive multitasking.

Although there were plans to upgrade the cooperative multitasking Mac OS to a preemptive model (and a preemptive API did exist in Mac OS 9, although in a very limited sense^[4] and rarely exploited), these were

NOTES

NOTES

abandoned in favor of Mac OS X, a hybrid of MacOS and the NextStep operating system, which is based on the Mach kernel and provides Unix-like preemptive multitasking.

OS/2 Warp, IBM's rewrite of an earlier IBM/Microsoft collaboration, OS/2, targeted at 386 systems, supported preemptive multitasking of native applications, and also permitted several different Windows sessions to be multitasked preemptively.

Non preemptive systems

Non preemptive multitasking is a style of computer multitasking in which the operating system never initiates a context switch from a running process to another process. Such systems are either statically scheduled, most often periodic systems, or exhibit some form of cooperative multitasking, in which case the computational tasks can self-interrupt and voluntarily give control to other tasks. When non preemptive is used, a process that receives such resources can not be interrupted until it is finished.

Cooperative multitasking is a type of multitasking in which the process currently controlling the CPU must offer control to other processes. It is called "cooperative" because all programs must cooperate for it to work. In contrast, preemptive multitasking interrupts applications and gives control to other processes outside of an application's control.

4.4 SYNCHRONIZATION

Usually the client sends requests to the server and receives results or error reports from the server. The client-server synchronization tools can extend this model. A client can be informed about *events* that occur on a server as a result of different client actions, such as data changes in a table or a procedure call.

Synchronization Principle

- The client can *register* certain event in order to be informed about the event occurrence. Occurrences before event registration are ignored.
- The client may then enter a *wait* state for any of his registered events occurrence (the waiting thread has a *state* = 10 in a list of clients).
- Another client may *invoke certain event* on the server.
- The clients waiting for this event to occur are then notified of its occurrence and their wait state is ended.
- The clients that have this event registered, but are not waiting for it, are informed of its occurrence if they decide to wait for it.

A single event occurrence therefore affects all clients that have this event registered.

Synchronization and Threads

When a client waits for an event, the waiting thread is blocked. If the client wants to make another actions when waiting, he has to run multiple threads. If the client wants to communicate with the server while waiting (e.g. make some database operations or stop the waiting state), the other thread has to connect to the server as well.

Connection to a server with different thread from the same process doesn't consume further licenses.

Event Labeling

Events are labeled by its *name* with maximum length of 31 characters. Registered events are also marked by an *event handler* number on the client side.

The name of an event is specified when the event is registered and when invoking the event. The event namespace is common for all clients. It's therefore recommended to use specific event names for each application (e.g. by using the application name as a prefix) if you don't want different application events on the same server to interact.

The event handler is assigned to a client when registering the event. The handler is returned after successful wait for the event. The handler is also used when unregistering the event.

Different clients have different handlers for events of the same name. If a client registers the same event repeatedly, it's assigned different handler.

Events and Transactions

If a client invokes an event, the event occurs after finishing the transaction. If the transaction is rolled back (explicitly or implicitly due to some error), then the event doesn't occur-invoking such an event will have no effect.

Event Parameters

You can set parameters as string characters when invoking an event. A client waiting for an event then receives these parameters.

You can set the required values of parameters when registering an event. Occurrences of an event that don't match this requirements are ignored.

There can be two types of parameter value requirements:

- either a specific value is required,
- or a string of parameters is required to begin with a specific value.

Parameter values comparison is not case-sensitive.

A client can register the same event multiple times with different (or even the same) parameter requirements. If an event with matching parameters occurs, each registration of this event will result in a successful wait stop.

NOTES

NOTES

In some cases, it is necessary to recognize an event occurrence in one application by a client working in another application. In another cases, such behavior is unwanted, e.g. when multiple instances of the same application are running independently on one server.

That is why a difference to *global* and *local* events was introduced. Local event occurrence is recognized only by a client that has set the same schema during the event commit as the client that has caused the event. Global event occurrence can be recognized by all clients that have registered it.

Multiple Events Occurrences

A wait for an event returns also the number of occurrences of that event. The wait comes to a stop on the first occurrence, however the number of occurrences (more than one) may be returned in these cases:

- the event was invoked multiple times during a single transaction;
- the event occurred before the client was waiting for it.

If an event occurs repeatedly with different parameter values, then these occurrences don't sum together. The individual parameter values can be obtained by the client by repeated calls of wait for this event.

If a client has an event registered, it is required to wait for this event with such parameter requirements, so it can be notified of this event. Otherwise event occurrences information would accumulate on the server.

4.5 SEMAPHORES

Semaphores are a programming construct designed by E.W. Dijkstra in the late 1960s. Dijkstra's model was the operation of railroads, consider a stretch of railroad in which there is a single track over which only one train at a time is allowed. Guarding this track is a semaphore. A train must wait before entering the single track until the semaphore is in a state that permits travel. When the train enters the track, the semaphore changes state to prevent other trains from entering the track. A train that is leaving this section of track must again change the state of the semaphore to allow another train to enter. In the computer version, a semaphore appears to be a simple integer. A process (or a thread) waits for permission to proceed by waiting for the integer to become 0. The signal if it proceeds signals that this by performing incrementing the integer by 1. When it is finished, the process changes the semaphore's value by subtracting one from it.

4.6 SEMAPHORE IMPLEMENTATION

Novell NetWare provides a simple semaphore facility that you can use to synchronize user access to network resources. This section describes NetWare's semaphore facility, suggests some possible applications, and illustrates how semaphores can be used to create a program that limits the number of users who can run it concurrently.

Semaphores

If you've had a few computer science classes, you've probably encountered semaphores. But unless you've written an operating system or a distributed application, chances are you've never used them. NetWare provides a simple implementation of semaphores that you can use to synchronize processes across the network or enforce exclusive access to a shared resource.

This section gives a brief introduction to semaphores, describes how semaphores are implemented in NetWare, then gives an example program that uses semaphores to implement concurrent-user licensing.

In 1965, Edsger W. Dijkstra proposed semaphores as a solution to the mutual exclusion problem. The mutual exclusion problem results when two processes share a resource.

For example, suppose your program has two processes (A and B), each of which is to increment a shared variable (X). The desired order of operations is:

Operation	Value of X
Initialize X.	5
A retrieves the value of X.	5
A increments X.	5
A stores the new value in X.	6
B retrieves the value of X.	6
B increments X.	6
B stores the new value in X.	7

Because the two processes run independently, the operations may occur in the following order:

Operation	Value of X
Initialize X.	5
A retrieves the value of X.	5

NOTES

NOTES

A increments X.	5
B retrieves the value of X.	5
A stores the new value in X.	6
B increments X.	6
B stores the new value in X.	6

Dijkstra's solution is to use a single variable, a semaphore, to control access to the shared resource. If the semaphore has a value greater than 0, the resource is available. If the semaphore has a value less than 0, processes are waiting for the shared resource. The two operations that operate on semaphores are called P (or wait) and V (or signal).

The operations get their name from the Dutch *proberen* (to test) and *verhogen* (to increment). Dijkstra defined the P operation on semaphore S as:

if $S > 0$ then $S = S - 1$ else (wait on S)

The V operation is:

if (processes are waiting on S) then (let one of the processes proceed)
else $S = S + 1$

The P and V operations must be implemented atomically, or the semaphore becomes just another shared resource with its own mutual exclusion problem.

Event Synchronization

One typical application of semaphores is in processing I/O requests. For example, a process may request data to be read from a file, then waits on a semaphore while another process services the request. The file read process signals the semaphore when the read is completed, and the first process can continue.

This interaction is called a *block/wakeup* protocol. Most implementations of semaphores (including NetWare's) allow for timeout after a specified period.

A similar type of process interaction is called a producer-consumer relationship. One process generates data and stores it in a shared buffer. The other process reads from the buffer. If the producer process operates faster than the consumer process, data will be lost.

If the consumer process runs faster than the producer process, the same data will be processed more than once. You can use semaphores to assure that the processes remain synchronized.

Counting Semaphores

Counting semaphores only take non-negative integer values, and are used to control access to a pool of identical resources. The example program

at the end of this DevNote, SYNC, uses counting semaphores to assure that no more than five users are allowed to run an application concurrently.

NetWare only has one type of semaphore, but they can operate as counting semaphores. See the description of SYNC for more information.

NOTES

NetWare's Synchronization APIs

NetWare has five API (application programming interface) functions that operate on semaphores:

- NW Open Semaphore.
- NW Close Semaphore.
- NW Wait On Semaphore.
- NW Signal Semaphore.
- NW Examine Semaphore.

Semaphore implementation in Unix

The System V specification of the UNIX operating system introduces semaphores among the other facilities for Inter-Process Communication (IPC). This implementation is a (rather messy) generalization of the above described scheme. Three features contribute to this unnecessary complication:

- A semaphore counter is not just a single value. It's rather a set of one or more values, and the exact number of values is specified upon semaphore creation.
- The creation of a semaphore is independent of its initialization. This is a fatal flaw, since we cannot atomically create a new semaphore set and initialize all the values in the set, a competing process may sneak in and access the semaphore after creation but before initialization, thus finding semagarbage.
- A semaphore must be explicitly destroyed when the shared resource is no longer necessary, since the system wide total number of semaphore sets is limited.

Moreover, semaphore counters are always nonnegative, a zero causing the waiting process to be put to sleep

A semaphore is created by the `semget(3)` system call, whose synopsis is:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int semget(key_t key, int nsems, int flag);
```

This system call returns an integer id for the semaphore (analogous to a file descriptor), or -1 in case of error.

NOTES

The key argument is analogous to a file name, and is used to identify exclusively a semaphore set (and other IPC objects as well) all over the system, while the semaphore id is like a file descriptor known only by one process and by its siblings that inherit it through variables after a `fork()`. Cooperating processes wishing to access the semaphore set must use the same key, much like a file name is used to store on disk shared data by cooperating processes. A key is implemented as an integer (whose size is specified by the `key_t` type), and a library function called `ftok()` is provided for the purpose of generating meaningful keys (see its manpage).

The value of `nsems` specifies the number of semaphores in newly created semaphore set.

The value of `flag` is a bit mask obtained oring constants that specify:

- The access permissions for the semaphore set, in the same way as for a file;
- The creation mode, which can be:
 - o `IPC_PRIVATE`: if this is specified, the value of the key is ignored, and a private semaphore set is created, that can be used by a process and its siblings through semaphore id inheritance. Since there's no key, other unrelated processes have no way to access the semaphore set.
 - o `IPC_CREAT`: if this is specified, and a semaphore with the given key does not exist, it is created, otherwise the call returns with `-1`, setting the appropriate `errno` value.

For example, to create one semaphore with key 333, readable and writable by an user and his group, you'd use a code fragment like:

```
int sid;
...
sid=semget((key_t)333, 1, SEM_R|SEM_W|(SEM_R>>3)|(SEM_W>>3)|IPC_CREAT);
if (sid<0)
{
    perror("Semaphore 333");
    exit(1);
}
...
```

Once a semaphore has been created, a process accesses it either directly via its id, if its a private semaphore and the id has been inherited from a creating ancestor, or by calling `semget()` to get its id in exchange for the right key, much like opening an already existing file. In this case the `IPC_PRIVATE` and `IPC_CREAT` bits are not set.

Semaphore initialization and flagging are performed respectively via the `semctl()` and `semop()` system calls.

4.7 MEMORY-MANAGEMENT ALLOCATION, SHARING AND MANIPULATING

NOTES

Due to the recent improvements in the price/performance characteristics of workstations and the networking capabilities, the client-server system architecture has become a target for database systems. A client-server management system provides the management of a database that resides on a client-server system. Data-access requests of the clients are handled by the database servers. The whole database is stored on the disks that can be accessed by only the servers, and in order to reduce disk accesses, copies of database items can be cached in the global memory which comprises the memories of all the computers connected to the system. Designing efficient global memory management algorithms helps the transactions experience less disk input/output (I/O) during their execution. A memory management algorithm for client-server aims to reduce disk I/O by increasing the portion of the database available in global memory.

Memory management is the act of managing computer memory. In its simpler forms, this involves providing ways to allocate portions of memory to programs at their request, and freeing it for reuse when no longer needed. The management of main memory is critical to the computer system.

Virtual memory systems separate the memory addresses used by a process from actual physical addresses, allowing separation of processes and increasing the effectively available amount of RAM using paging or swapping to secondary storage. The quality of the virtual memory manager can have a big impact on overall system performance.

Garbage collection is the automated allocation and deallocation of computer memory resources for a program. This is generally implemented at the programming language level and is in opposition to manual memory management, the explicit allocation and deallocation of computer memory resources. Region-based memory management is an efficient variant of explicit memory management that can deallocate large groups of objects simultaneously.

The memory management subsystem is one of the most important parts of the operating system. Since the early days of computing, there has been a need for more memory than exists physically in a system. Strategies have been developed to overcome this limitation and the most successful of these is virtual memory. Virtual memory makes the system appear to have more memory than it actually has by sharing it between competing processes as they need it.

Virtual memory does more than just make your computer's memory go further. The memory management subsystem provides:

Large Address Spaces

The operating system makes the system appear as if it has a larger amount of memory than it actually has. The virtual memory can be many times larger than the physical memory in the system.

Protection

NOTES

Each process in the system has its own virtual address space. These virtual address spaces are completely separate from each other and so a process running one application cannot affect another. Also, the hardware virtual memory mechanisms allow areas of memory to be protected against writing. This protects code and data from being overwritten by rogue applications.

Memory Mapping

Memory mapping is used to map image and data files into a processes address space. In memory mapping, the contents of a file are linked directly into the virtual address space of a process.

Fair Physical Memory Allocation

The memory management subsystem allows each running process in the system a fair share of the physical memory of the system.

Shared Virtual Memory

Although virtual memory allows processes to have separate (virtual) address spaces, there are times when you need processes to share memory. For example there could be several processes in the system running the **bash** command shell. Rather than have several copies of **bash**, one in each processes virtual address space, it is better to have only one copy in physical memory and all of the processes running **bash** share it. Dynamic libraries are another common example of executing code shared between several processes.

Shared memory can also be used as an Inter Process Communication (IPC) mechanism, with two or more processes exchanging information via memory common to all of them. Linux supports the UnixTM System V shared memory IPC.

An Abstract Model of Virtual Memory




Fig. 4.1 : Abstract model of Virtual to Physical address mapping

Before considering the methods that Linux uses to support virtual memory it is useful to consider an abstract model that is not cluttered by too much detail.

As the processor executes a program it reads an instruction from memory and decodes it. In decoding the instruction it may need to fetch or store the contents of a location in memory. The processor then executes the instruction and moves onto the next instruction in the program. In this way the processor is always accessing memory either to fetch instructions or to fetch and store data.

In a virtual memory system all of these addresses are virtual addresses and not physical addresses. These virtual addresses are converted into physical addresses by the processor based on information held in a set of tables maintained by the operating system.

To make this translation easier, virtual and physical memory are divided into handy sized chunks called *pages*. These pages are all the same size, they need not be but if they were not, the system would be very hard to administer. Linux on Alpha AXP systems uses 8 Kbyte pages and on Intel x86 systems it uses 4 Kbyte pages. Each of these pages is given a unique number; the page frame number (PFN).

In this paged model, a virtual address is composed of two parts; an offset and a virtual page frame number. If the page size is 4 Kbytes, bits 11:0 of the virtual address contain the offset and bits 12 and above are the virtual page frame number. Each time the processor encounters a virtual address it must extract the offset and the virtual page frame number. The processor must translate the virtual page frame number into a physical one and then access the location at the correct offset into that physical page. To do this the processor uses *page tables*.

Figure 4.1 shows the virtual address spaces of two processes, process X and process Y, each with their own page tables. These page tables map each processes virtual pages into physical pages in memory. This shows that process X's virtual page frame number 0 is mapped into memory in physical page frame number 1 and that process Y's virtual page frame number 1 is mapped into physical page frame number 4. Each entry in the theoretical page table contains the following information:

- Valid flag. This indicates if this page table entry is valid.
- The physical page frame number that this entry is describing.
- Access control information. This describes how the page may be used. Can it be written to? Does it contain executable code?

The page table is accessed using the virtual page frame number as an offset. Virtual page frame 5 would be the 6th element of the table (0 is the first element).

To translate a virtual address into a physical one, the processor must first work out the virtual addresses page frame number and the

NOTES

NOTES

offset within that virtual page. By making the page size a power of 2 this can be easily done by masking and shifting. Looking again at Fig. 4.1 and assuming a page size of $0x2000$ bytes (which is decimal 8192) and an address of $0x2194$ in process Y's virtual address space then the processor would translate that address into offset $0x194$ into virtual page frame number 1.

The processor uses the virtual page frame number as an index into the processes page table to retrieve its page table entry. If the page table entry at that offset is valid, the processor takes the physical page frame number from this entry. If the entry is invalid, the process has accessed a non-existent area of its virtual memory. In this case, the processor cannot resolve the address and must pass control to the operating system so that it can fix things up.

Just how the processor notifies the operating system that the correct process has attempted to access a virtual address for which there is no valid translation is specific to the processor. However the processor delivers it, this is known as a *page fault* and the operating system is notified of the faulting virtual address and the reason for the page fault.

Assuming that this is a valid page table entry, the processor takes that physical page frame number and multiplies it by the page size to get the address of the base of the page in physical memory. Finally, the processor adds in the offset to the instruction or data that it needs.

Using the above example again, process Y's virtual page frame number 1 is mapped to physical page frame number 4 which starts at $0x8000$ ($4 \times 0x2000$). Adding in the $0x194$ byte offset gives us a final physical address of $0x8194$.

By mapping virtual to physical addresses this way, the virtual memory can be mapped into the system's physical pages in any order. For example, in Fig. 4.1 process X's virtual page frame number 0 is mapped to physical page frame number 1 whereas virtual page frame number 7 is mapped to physical page frame number 0 even though it is higher in virtual memory than virtual page frame number 0. This demonstrates an interesting byproduct of virtual memory; the pages of virtual memory do not have to be present in physical memory in any particular order.

SUMMARY

1. Scheduling disciplines are algorithms used for distributing resources among parties which simultaneously and asynchronously request them.
2. Preemptive scheduling allows a process to be interrupted in the midst of its execution, taking the CPU away and allocating it to another process.
3. The term preemptive multitasking is used to distinguish a multitasking operating system, which permits preemption of tasks, from a cooperative multitasking system wherein processes or tasks must be explicitly programmed to yield when they do not need system resources.

4. The client can register certain event in order to be informed about the event occurrence. Occurrences before event registration are ignored.
5. When the train enters the track, the semaphore changes state to prevent other trains from entering the track. A train that is leaving this section of track must again change the state of the semaphore to allow another train to enter.
6. Memory management is the act of managing computer memory. In its simpler forms, this involves providing ways to allocate portions of memory to programs at their request, and freeing it for reuse when no longer needed.

NOTES

REVIEW QUESTIONS

1. What is Scheduling and how we can implement this in client server systems?
2. What is a scheduler explain with reference to client server process communications?
3. What is a preemptive operating system explain with example?
4. What are non-preemptive systems explain with examples?
5. Why we need synchronization and how we can implement this with using semaphores?
6. Explain semaphore implementation UNIX.
7. Define Memory-management and Allocation with sharing and manipulating the processes in client server systems.

NOTES

CHAPTER

5

CLIENT SERVER APPLICATIONS DESIGN AND DEVELOPMENT

STRUCTURE

- 5.1 Client server computing with Oracle-Overview of DBMS
- 5.2 Client server relationships
- 5.3 Oracle and client server computing using SQL
- 5.4 The oracle tools and design aids
- 5.5 SQL windows and Power Builder
 - *Summary*
 - *Review Questions*

5.1 CLIENT SERVER COMPUTING WITH ORACLE-OVERVIEW OF DBMS

An Oracle database is a collection of data treated as a unit. The purpose of a database is to store and retrieve related information. A database server is the key to solving the problems of information management. In general, a server reliably manages a large amount of data in a multiuser environment so that many users can concurrently access the same data. All this is accomplished while delivering high performance. A database server also prevents unauthorized access and provides efficient solutions for failure recovery.

Oracle Database is the first database designed for enterprise grid computing, the most flexible and cost effective way to manage information and applications. Enterprise grid computing creates large pools of industry-standard, modular storage and servers. With this architecture, each new

system can be rapidly provisioned from the pool of components. There is no need for peak workloads, because capacity can be easily added or reallocated from the resource pools as needed.

The database has logical structures and physical structures. Because the physical and logical structures are separate, the physical storage of data can be managed without affecting the access to logical storage structures.

Overview of Application Architecture

There are two common ways to architect a database: client/server or multitier. As internet computing becomes more prevalent in computing environments, many database management systems are moving to a multitier environment.

Client Server Architecture

Multiprocessing uses more than one processor for a set of related jobs. Distributed processing reduces the load on a single processor by allowing different processors to concentrate on a subset of related tasks, thus improving the performance and capabilities of the system as a whole.

An Oracle database system can easily take advantage of distributed processing by using its client server architecture. In this architecture, the database system is divided into two parts: a front-end or a client, and a back-end or a server.

The Client

The client is a database application that initiates a request for an operation to be performed on the database server. It requests, processes, and presents data managed by the server. The client workstation can be optimized for its job. For example, it might not need large disk capacity, or it might benefit from graphic capabilities.

Often, the client runs on a different computer than the database server, generally on a PC. Many clients can simultaneously run against one server.

The Server

The server runs Oracle software and handles the functions required for concurrent, shared data access. The server receives and processes the SQL and PL/SQL statements that originate from client applications. The computer that manages the server can be optimized for its duties. For example, it can have large disk capacity and fast processors.

Multitier Architecture, Application Servers

A multitier architecture has the following components:

- A client or initiator process that starts an operation
 - One or more application servers that perform parts of the operation.
- An application server provides access to the data for the client and performs some of the query processing, thus removing some

NOTES

NOTES

of the load from the database server. It can serve as an interface between clients and multiple database servers, including providing an additional level of security.

- An end or database server that stores most of the data used in the operation.

This architecture enables use of an application server to do the following:


- Validate the credentials of a client, such as a Web browser.
- Connect to an Oracle database server.
- Perform the requested operation on behalf of the client.

If proxy authentication is being used, then the identity of the client is maintained throughout all tiers of the connection.

5.2 CLIENT SERVER RELATIONSHIPS

The “client-server relationship” is a very common term when it comes to file transfer, but what *exactly* is a “client” and what is a “server”? The networking of a “client-server relationship” involves a distributed application architecture that divides tasks between service *providers* known as servers and a service *requester*, or a client. Clients and servers generally operate over computer networks, such as office LANs or even the Internet.

The server is the more dominant machine in the relationship, being a high-performance host that runs one (or more than one) server applications.



The server applications are what the client “talks to”. Applications on the client system initiate communication sessions with servers.

Simply put, “client-server” describes the relationship between two computer programs where the client makes request to the server. This relationship is seen in such tasks as sending emails and accessing the

NOTES

web. Many business applications currently use the client-server model as well as the main Internet application protocols (HTTP, SMTP, telnet etc).

The simplest type of client-server uses only two systems, a client and a server. This may also be referred to as "two-tier" and it allows for the sharing of files. Some examples of clients in this relationship would be web browsers, email clients, and online chat clients. Servers would include web servers, ftp servers, application servers, mail servers, file servers, and most web services.

Client-server architecture enables the responsibilities of a system to be distributed amongst several individual computers. These individual computers are accessed only through a network so it is easier to maintain and repair or even relocate a server without affecting the clients. Another benefit exists in that data is stored on servers which generally have greater security than clients. Servers can also control user access and resources. Client-server technologies are abundant in the world today, designed to ensure safety and security. One example of an innovator in this field is Ipswitch—a software company which has been offering FTP software since 1991. As the field of file transfer continues to grow, so will the need for safe servers and happy clients who together form a healthy, networked relationship.

5.3 ORACLE AND CLIENT SERVER COMPUTING

In client/server architecture, clients—programs that represent users who need services—and servers—programs that provide services—are separate logical objects that communicate over a network to perform tasks together. A client makes a request for a service and receives a reply to that request. A server receives and processes a request, and sends back the required response.

Characteristics of Client/Server Architecture

The client/server architecture has the following characteristics:

- Asymmetrical protocols—a many-to-one relationship between clients and a server. Clients always initiate a dialog by requesting a service. Servers wait passively for requests from clients.
- Encapsulation of services—the server is a specialist, when given a message requesting a service, it determines how to get the job done. Servers can be upgraded without affecting clients as long as the published message interface used by both is unchanged.
- Integrity—the code and data for a server are centrally maintained, which results in inexpensive maintenance and the protection of shared data integrity. At the same time, clients remain personal and independent.

NOTES

- Location transparency—the server is a process that can reside on the same machine as a client process or on a different machine across a network. Client/server software usually hides the location of a server from clients by redirecting service requests. Clients should not have to be aware of the location of servers.
- Namespace transparency—clients should be able to use the same naming conventions (and namespace) to locate any server on the network.
- Message-based exchanges—clients and servers are loosely-coupled processes that can exchange service requests and replies using messages.
- Modular, extensible design—the modular design of a client/server application enables that application to be fault-tolerant. In a fault-tolerant system, failures may occur without causing a shutdown of the entire application. In a fault-tolerant client/server application, one or more servers may fail without stopping the whole system as long as the services offered on the failed servers are available on servers that are still active. Another advantage of modularity is that a client/server application can respond automatically to increasing or decreasing system loads by adding or shutting down one or more services or servers.
- Platform independence—the ideal client/server software is independent of hardware or operating system platforms, allowing the mixing of client and server platforms. Clients and servers can be deployed on different hardware using different operating systems, optimizing the type of work each performs.
- Reusable code—service programs can be used on multiple servers.
- Scalability—client/server systems can be scaled horizontally or vertically. Horizontal scaling means adding or removing client workstations with only a slight performance impact. Vertical scaling means migrating to a larger and faster server machine or adding server machines.
- Separation of client/server functionality—client/server is a relationship between processes running on the same or separate machines. A server process is a provider of services. A client is a consumer of services. Client/server provides a clean separation of functions.
- Shared resources—one server can provide services for many clients at the same time, and regulate their access to shared resources.

Differences Between 2-Tier and 3-Tier Client/Server Architectures

Every client/server application contains three functional units:

- Presentation logic or user interface (for example, ATM machines).
- Business logic (for example, software that enables a customer to request an account balance).

- Data (for example, records of customer accounts).

NOTES

These functional units can be part of the client program or part of the one or more server programs in your application. Which of the many possible variations you choose depends on how you split the application and which middleware you use to communicate between the tiers, as illustrated in the following figure.

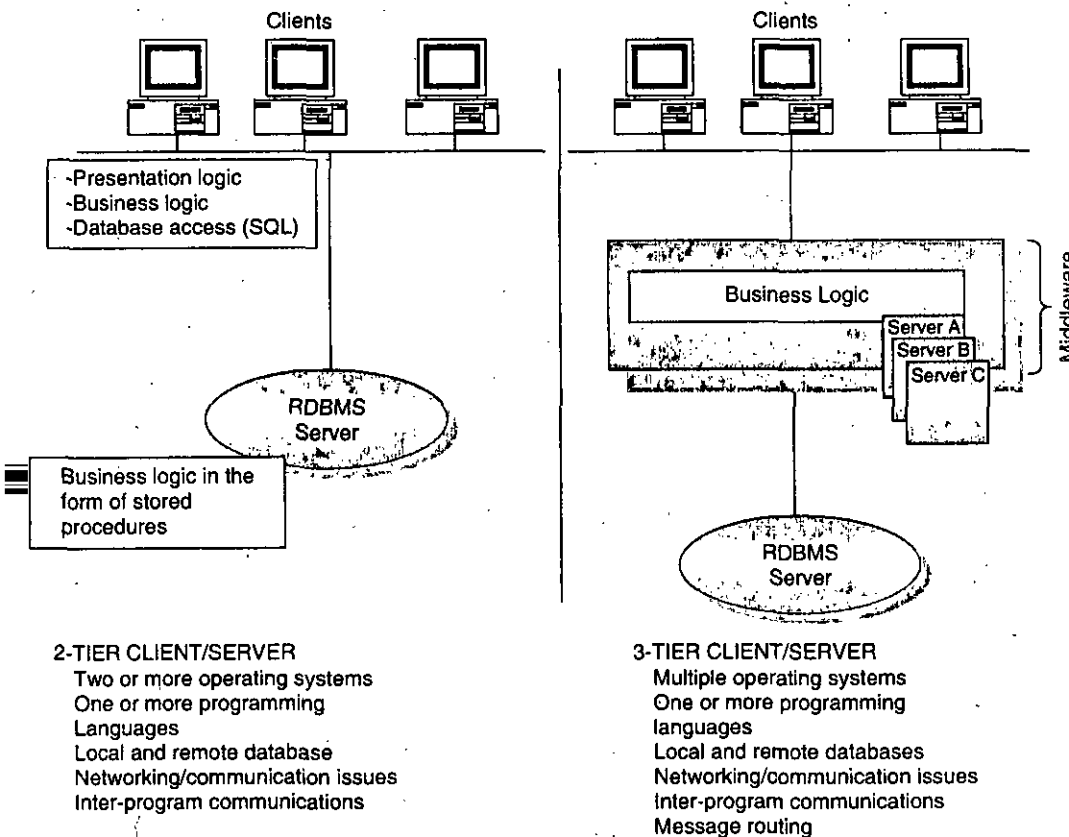


Fig. 5.1 : 2-Tier and 3-Tier Client/Server Models

In 2-tier client/server applications, the business logic is buried inside the user interface on the client or within the database on the server in the form of stored procedures. Alternatively, the business logic can be divided between the client and server. File servers and database servers with stored procedures are examples of 2-tier architecture.

In 3-tier client/server applications, the business logic resides in the middle tier, separate from the data and user interface. In this way, processes can be managed and deployed separately from the user interface and the database. Also, 3-tier systems can integrate data from multiple sources.

Client Server Variations to Suit Your Needs

Client/server architecture can accommodate the needs of each of the following situations:

NOTES

- Small shops and laptops—the client, the middleware software, and most of the business services operate on the same machine. Oracle recommends this approach for one-person businesses such as a dentist's office, a home office, and a business traveler who frequently works on a laptop computer.
- Small businesses and corporate departments—a LAN-based single-server application is required. Users of this type of application include small businesses, such as a medical practice with several doctors, a multi-department corporation, or a bank with several branch offices. In this type of application, multiple clients talk to a local server. Administration is simple security is implemented at the machine level and failures are detected easily.
- Large enterprises—multiple servers that offer diverse functionality are required. Multiple servers can reside on corporate networks, intranets, and the Internet, all of which are highly scalable. Servers can be partitioned by function, resources, or databases, and can be replicated for increased fault tolerance or enhanced performance. This model provides a great amount of power and flexibility. How well you architect your application is critical to this client/server model. You may need to partition work among servers, or design servers to delegate work to other servers.

How the Oracle Fits into the Client Server Model

The Oracle system fits into the middle of the client/server model. In a Oracle Tuxedo application, clients log in and request services offered by an application. The Oracle Tuxedo system offers these services through a transparent bulletin board. The bulletin board provides a *global directory advertising service*.

For example, in the following sample banking application, the bulletin board advertises deposit, withdrawal, and inquiry services. The Oracle Tuxedo system then finds a server at the appropriate branch or district office that can provide the requested services.

The sample banking application shows the primary building blocks of a Oracle Tuxedo application:

- Clients—programs that collect input from users, send requests through the Oracle Tuxedo system to servers, and deliver server replies to users.
- Servers—programs that encapsulate the business logic into a set of services that define the application.
- Middleware—comprises all the distributed software needed to support interactions between clients and servers. It is the medium that enables a client to obtain a service from a server. Middleware includes (1) API functions used by the client—to issue requests and receive replies—and the server—to issue replies—and (2) messaging

paradigms used to transmit client requests and server responses over a network. Middleware does not include the client user interface, the application logic, or the services provided by the servers.

NOTES

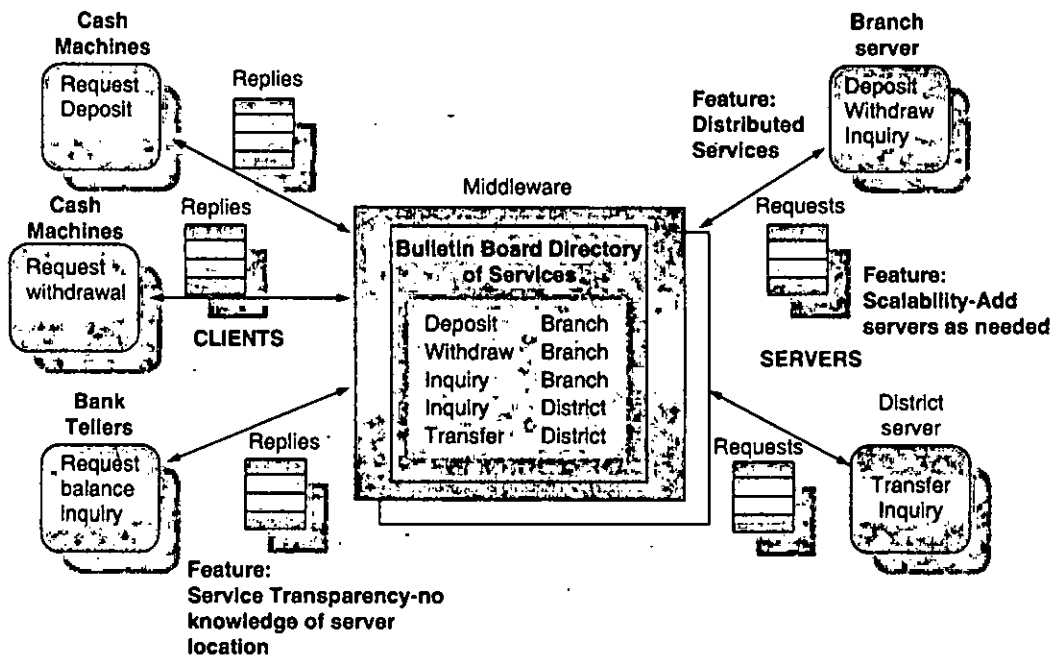


Fig. 5.2 : Clients and Servers in a Sample Banking Application

In the sample Oracle Tuxedo banking application, clients (cash machines and tellers) make requests, and servers (at branch and district offices) provide services and responses. For example, a customer may use a cash machine to find out how much money is available in his personal checking account. The cash machine (a client) calls the server to get the balance. The server receives the request, retrieves the balance, and sends the information to the cash machine.

What Is a Oracle Client

A client is a program that collects a request from a user and passes that request to a server capable of fulfilling it. It can reside on a PC or workstation as part of the front end of an application. It can also be embedded in software that reads a communication device such as an ATM machine from which data is collected and formatted before being processed by Oracle Tuxedo servers.

To be a client, a program must be able to invoke the Oracle Tuxedo libraries of functions and procedures known collectively as the Application-to-Transaction-Monitor Interface, or ATMI. The ATMI is supported in several language bindings.

A client joins a Tuxedo application by calling the ATMI client initialization routine. Once it has joined an application, a client can define transaction boundaries and call ATMI functions that enable it to communicate with

NOTES

other programs in the application. The client leaves the application by issuing an ATMI termination function. By joining an application only when necessary and leaving it once the appropriate task is complete, a client frees Oracle Tuxedo system resources for use by other clients and servers.

When building a distributed application, you must determine how information is gathered and presented to your business for processing. You have complete control over where and when to call ATMI functions, depending upon your business logic and rules. Your program can join one Oracle Tuxedo application, perform some tasks and leave, and then join a different Oracle Tuxedo application to perform another task. If you are using a multicontexted application, your client can perform tasks in more than one application without leaving any of them.

What Is a Oracle Server

A Oracle Tuxedo server is a process that oversees a set of services, dispatching them automatically for clients that request them. A service, in turn, is a function within a server program that performs a particular task needed by a business. A bank, for example, might have one service that accepts deposits and another that reports account balances. A server at this bank might receive requests from clients for both services. The server is responsible for dispatching each request to the appropriate service.

Service functions implement business logic through calls to database interfaces such as SQL and, possibly, calls to the ATMI to access additional services, queues, and other resources. The servers on which these services reside then reply to the clients or send the client requests to a new service.

Application Processing Services Provided by the Oracle

The Oracle Tuxedo system provides services that enable application developers to implement the following functionality in their applications:

- Data compression.
- Data-dependent routing.
- Data encoding.
- Data encryption.
- Data marshalling.
- Load balancing.
- Message prioritization.
- Service and event naming.

For descriptions of the Tuxedo application processing services, see Oracle Tuxedo ATMI Architecture.

Administrative Services Provided by the Oracle

The Oracle Tuxedo system provides services that enable application administrators to perform the following administrative tasks:

- Startup and shutdown of an application.

- Centralized application configuration.
- Distributed application management.
- Dynamic application reconfiguration.
- Workstation management.
- Security management.
- Transaction management.
- Message queuing management.
- Event management.

NOTES

5.4 THE ORACLE TOOLS AND DESIGN AIDS

Oracle Corporation's reputation as a database company is firmly established in its full-featured, high-performance RDBMS server. With the database as the cornerstone of its product line, Oracle has evolved into more than just a database company, complementing its RDBMS server with a rich offering of well-integrated products that are designed specifically for distributed processing and client/server applications. As Oracle's database server has evolved to support large-scale enterprise systems for transaction processing and decision support, so too have its other products, to the extent that Oracle can provide a complete solution for client/server application development and deployment. This chapter presents an overview of client/server database systems and the Oracle product architectures that support their implementation.

An Overview of Client/Server Computing

The premise of client/server computing is to distribute the execution of a task among multiple processors in a network. Each processor is dedicated to a specific, focused set of subtasks that it performs best, and the end result is increased overall efficiency and effectiveness of the system as a whole. Splitting the execution of tasks between processors is done through a protocol of service requests; one processor, the *client*, requests a service from another processor, the *server*. The most prevalent implementation of client/server processing involves separating the user interface portion of an application from the data access portion.

On the client, or *front end*, of the typical client/server configuration is a user workstation operating with a Graphical User Interface (GUI) platform, usually Microsoft Windows, Macintosh, or Motif. At the *back end* of the configuration is a database server, often managed by a UNIX, Netware, Windows NT, or VMS operating system.

Client/server architecture also takes the form of a *server-to-server* configuration. In this arrangement, one server plays the role of a client, requesting database services from another server. Multiple database servers can look like a single logical database, providing transparent access to data that is spread around the network.

NOTES

Designing an efficient client/server application is somewhat of a balancing act, the goal of which is to evenly distribute execution of tasks among processors while making optimal use of available resources. Given the increased complexity and processing power required to manage a graphical user interface (GUI) and the increased demands for throughput on database servers and networks, achieving the proper distribution of tasks is challenging. Client/server systems are inherently more difficult to develop and manage than traditional host-based application systems because of the following challenges:

- The components of a client/server system are distributed across more varied types of processors. There are many more software components that manage client, network, and server functions, as well as an array of infrastructure layers, all of which must be in place and configured to be compatible with each other.
- The complexity of GUI applications far outweighs that of their character-based predecessors. GUIs are capable of presenting much more information to the user and providing many additional navigation paths to elements of the interface.
- Troubleshooting performance problems and errors is more difficult because of the increased number of components and layers in the system.

Databases in a Client/Server Architecture

Client/server technologies have changed the look and architecture of application systems in two ways. Not only has the supporting hardware architecture undergone substantial changes, but there have also been significant changes in the approach to designing the application logic of the system.

Prior to the advent of client/server technology, most Oracle applications ran on a single node. Typically, a character-based SQL*Forms application would access a database instance on the same machine with the application and the RDBMS competing for the same CPU and memory resources. Not only was the system responsible for supporting all the database processing, but it was also responsible for executing the application logic. In addition, the system was burdened with all the I/O processing for each terminal on the system; each keystroke and display attribute was controlled by the same processor that processed database requests and application logic.

Client/server systems change this architecture considerably by splitting all of the interface management and much of the application processing from the host system processor and distributing it to the client processor.

Combined with the advances in hardware infrastructure, the increased capabilities of RDBMS servers have also contributed to changes in the application architecture. Prior to the release of Oracle7, Oracle's RDBMS was less sophisticated in its capability to support the processing logic necessary to maintain the integrity of data in the database. For example,

NOTES

primary and foreign key checking and enforcement was performed by the application. As a result, the database was highly reliant on application code for enforcement of business rules and integrity, making application code bulkier and more complex. Figure 5.1 illustrates the differences between traditional host-based applications and client/server applications. Client/server database applications can take advantage of the Oracle 7 server features for implementation of some of the application logic.

The pre-Oracle 7 databases provided little more than data-type checking at the kernel level; but with Oracle 7, much of the application logic processing can be performed by the database kernel. Oracle 7 contains features such as stored procedures, integrity constraint enforcement, user-defined functions, and database triggers, all of which enable the application to store more of its business rules (or semantics of the data model) at the database level. As a result, the application is freed to do more sophisticated, complex processing tasks such as GUI interface management and integration to other client-based productivity tools. As shown in Fig. 5.1, the database is much more robust; no longer is it reliant on application code to maintain its integrity.

Oracle and Client/Server Computing

Oracle Corporation has been a leader in introducing advanced client/server database technologies, directing its product development specifically to support the design, implementation, and management of client/server database systems. Oracle has designed products to support each of the three primary components of a client/server architecture:

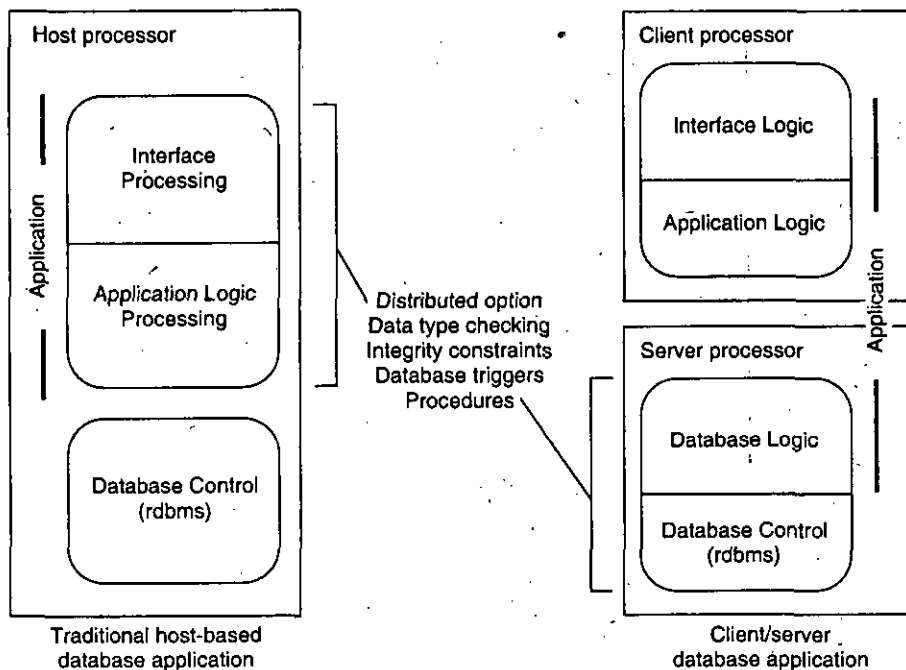


Fig. 5.3 : Oracle's client/server model

NOTES

- A full-featured, high-performance RDBMS server, scaleable from laptops to mainframes.
- Client development and run-time products that support multiple GUI environments.
- Database connectivity middleware that provides efficient and secure communications over a wide variety of network protocols.

Oracle's product offerings in each area are highly scaleable, providing complete client/server solutions for application environments ranging from small workgroup to global enterprise-wide environments.

The following sections describe several of the primary components of Oracle's client/server architecture, the Oracle7 RDBMS kernel, SQL*Net, and GUI development tools, all designed specifically for client/server application systems.

The Oracle7 RDBMS Server

The Oracle7 server is a full-featured RDBMS that is ideally suited to support sophisticated client/server environments. Many features of the Oracle7 internal architecture are designed to provide high availability, maximum throughput, security, and efficient use of its host's resources. Although all these features are important architecturally for a database server, Oracle7 also contains the following language-based features that accelerate development and improve the performance of server-side application components.

PL/SQL language—A major component of the Oracle7 server is its PL/SQL processing engine (the PL stands for Procedural Language). PL/SQL is Oracle's fourth generation language that incorporates structured procedural language elements with the SQL language. PL/SQL is designed specifically for client/server processing in that it enables a PL/SQL program block containing application logic as well as SQL statements to be submitted to the server with a single request.

By using PL/SQL, you can significantly reduce the amount of processing required by the client portion of an application and the network traffic required to execute the logic. For example, you might want to execute different sets of SQL statements based on the results of a query. The query, the subsequent SQL statements, and the conditional logic to execute them can all be incorporated into one PL/SQL block and submitted to the server in one network trip.

Not only can PL/SQL be processed by the Oracle7 server, but it can also be processed by SQL* Forms and Oracle Forms. PL/SQL is used extensively by these tools for client-based procedures and event trigger routines. In a client/server environment, PL/SQL is extremely flexible because the language used by client is interchangeable with that used by the server. Some extensions in the client language syntax allow for control of interface components, reference to form objects, and navigation. Fig. 5.3 shows the difference

between submitting an “anonymous” PL/SQL block and calling a stored procedure.

NOTES

- **Stored procedures**—Although version 6 of Oracle supported server-based PL/SQL, Oracle7 provides the capability to store PL/SQL blocks as database objects in the form of stored procedures, functions, and database packages. Now, portions of the application logic, especially those requiring database access, can reside where they are processed—on the server. Using stored procedures significantly increases the efficiency of a client/server system for several reasons:
- Calling a stored procedure from a client application generates minimal network traffic. Rather than the application submitting an entire PL/SQL program block from the client, all that is required is a single call to the procedure or function with an optional parameter list.
- Stored procedures provide a convenient and effective security mechanism. One of the characteristics of stored PL/SQL is that it always executes with the privilege domain of the procedure owner. This enables non-privileged users to have controlled access (through the procedure code) of privileged objects. This feature usually serves to reduce the amount of grant administration that the DBA must do.
- Both the compiled and textual form of stored procedures are maintained in the database. Because the compiled form of the procedure is available and readily executable, the need to parse and compile the PL/SQL at run time is alleviated.
- **Database triggers**—Database triggers resemble stored procedures in that they are database-resident PL/SQL blocks; the difference between the two is that triggers are fired automatically by the RDBMS kernel in response to a commit time event (such as an insert, update, or delete operation). You can use triggers to enforce complex integrity checking, perform complex auditing and security functions, and implement application alerts and monitors. Like stored procedures, database triggers greatly reduce the amount of code and processing that is necessary in the client portion of an application.
- Oracle7's implementation of database triggers is slightly different from that of other vendors. Although most databases support statement-level triggers, Oracle7 also includes functionality to fire triggers at the row-level. Consider an UPDTAE statement that affects values in a set of 100 rows. The kernel would fire a statement-level trigger once—for the UPDATE statement (either before and/or after the statement executes). Row-level triggers, on the other hand, are fired by the kernel for each row that the statement

NOTES

affects—in this case, 100 times. Oracle7 enables statement-level and row-level triggers to be used in conjunction with one another.

- **Declarative integrity**—When you define a table in Oracle7, you might include integrity constraints as part of your table definition. Constraints are enforced by the server whenever records are inserted, updated, or deleted. In addition to using referential integrity constraints that enforce primary and foreign key relationships, you can also define your own constraints to control the value domains of individual columns within a table.
- **Server-enforced integrity** reduces some of the code required for validation by the client and also increases the robustness of the business model defined within the database. With constraints, you can often improve performance and provide the flexibility to support multiple front-end interfaces.
- **User-defined functions**—You'll also find PL/SQL blocks in user-defined functions. User-defined functions are similar to stored procedures and also reduce the amount of application code in the client portion of an application. Not only can you call these functions from PL/SQL, but you can also use them to extend the set of standard Oracle SQL functions. You can place user-defined functions in SQL statements just as you would any other Oracle SQL function.

Designing your Oracle application to make use of these server-based features not only improves the performance of a client/server system but also makes the task of developing and deploying an application easier.

Networking Products

If you're developing an Oracle-based client/server system, you'll probably use Oracle's database networking software to implement connectivity between the nodes in the network. Oracle offers a variety of products and tools that simplify the task of connecting client applications to database servers in a network.

- **SQL*Net**—SQL*Net is database messaging software that provides optimal, reliable database messaging over every popular network protocol. SQL*Net is designed to provide server location transparency to any node within an application network and uses components that reside on both the client and server sides of an application. In addition to providing connectivity between workstations and servers in a client/server environment, servers also use SQL*Net to communicate with other servers for distributed transactions, remote procedure calls, and table replication. Servers reference other servers using *database links* that define the names of remote databases, the network nodes where the databases are serviced, and the network protocol used to access the remote nodes. Database links simplify distributed processing by providing transparent access

to remote objects such as tables and procedures, enabling an application to reference them just as if they were resident in the application's local database.

NOTES

- **Oracle Names**—With release 7.1 of Oracle 7, you can make available database link and network node information to all the nodes in a network using the Oracle Names common global dictionary. This feature is particularly useful for large application networks encompassing multiple locations to simplify the administration of database link and network information.
- **Multi-Protocol Interchange**—Whereas Version 1 of SQL*Net supports connectivity between nodes in a single network protocol, Version 2 of SQL*Net enables database communication between nodes in different network communities running different network protocols. The Multi-Protocol Interchange (MPI) provides a communications bridge over heterogeneous protocols by translating SQL*Net messages from one protocol to another. For example, a client workstation in a Token-Ring LAN can transparently access a server in a DECnet or TCP/IP network, insulating the application from the complexities of the underlying network infrastructure. In addition to providing multi-protocol communications, the MPI also provides cost-based message routing functions and uses alternate routes in the event that least-cost paths of a network are unavailable.
- **Oracle Network Manager**—The complex task of configuring and managing a distributed database network topology is made easier with the Network Manager, a GUI-based administrative interface for SQL*Net. The Network Manager is used not only to manage the Oracle Names dictionary but also to generate configuration files for client and server-side SQL*Net components and define connection routes for Multi-Protocol Interchange nodes.

Client/Server Development Tools

In addition to its server and networking products, Oracle includes a variety of client-side GUI offerings that complete its integrated client/server architecture. These product suites include full-featured Computer Assisted Software Engineering (CASE) tools, object-oriented development environments, and run-time components that are capable of operating with the Oracle 7 server as well as other SQL databases.

- **Designer/2000**—For developing sophisticated Oracle client/server applications, the Designer/2000 CASE environment provides a comprehensive repository and powerful tool set that enables you to systematically analyze, model, design, and generate both client and server components of an application. Designer/2000's repository is similar in most respects to its predecessor's, Oracle*CASE; however, its user interface and functionality are significantly enhanced,

and it supports graphical business process reengineering as well as Oracle7 server and Developer/2000 features.

NOTES

Designer/2000 is a complex tool with many features. Once you've mastered it, you'll find that the task of developing sophisticated applications from start to finish is faster and much more efficient than it is with traditional development methods. Information collected through its graphical diagrammers for data and application modeling is used to generate sophisticated, bug-free data definition language and application code. Designer/2000 generates full-featured, full-functioning applications for Developer/2000 tools, complete with menus, security, transaction control, and extensibility for OLE containers. Application generators are also in the works for Visual Basic and Power Objects.

- **Developer/2000**—The Developer/2000 suite packages Oracle Forms, Oracle Reports, Oracle Graphics, and Oracle Book into a single integrated development environment. You can build your applications using just these tools or use them in conjunction with Designer/2000 to produce generated forms and reports.

As with their character-based predecessors, SQL*Forms and SQL*ReportWriter, the Developer/2000 tools use PL/SQL as their underlying scripting language. Applications developed on one type of workstation platform can be deployed on other platforms such as Microsoft Windows, Macintosh, and Motif.

- **Power Objects**—In addition to the Designer/2000 and Developer/2000 tools, the Oracle GUI product suite includes another object-oriented, GUI application environment that is designed to compete with the likes of PowerBuilder and Visual Basic. Power Objects provides a rapid application development environment with lots of drag-and-drop features plus automatic database transaction management. Instead of using the PL/SQL language used by Developer/2000, Power Objects uses a Basic-style scripting language that is similar in most respects to Microsoft Visual Basic and is well suited for small to mid-sized applications.

Oracle Directions

Oracle Corporation continues to be a forward-thinking leader in the technology sector, introducing new technologies that provide extended functionality and better management, development capability, connectivity, and performance for scaleable client/server database systems. Some of the new technologies introduced as of this writing include the following:

- **Wireless client/server**—Oracle has recently introduced connectivity and messaging technologies that support remote client/server computing. This technology, which operates over cellular data networks, is particularly useful for laptop and palmtop users. A slightly more advanced form of client/server computing, *Oracle Mobile Agents*

uses a *client-agent-server* architecture that enables a client to work offline and connect periodically to the network to submit requests and receive results from the server. The *agent* component of this architecture functions on behalf of the client in its absence from the network.

- **Internet/World Wide Web interface**—Oracle's Web Interface Kit is used to integrate World Wide Web servers to Oracle7 databases. The kit provides utilities to create Web pages and interface Oracle7 data to them, enabling both storage and retrieval of Oracle data by Web users.
- **Multimedia server**—As applications use more different types of data, especially multimedia data, database server technologies will also improve to manage them. In future releases of the Oracle server, you'll see increased support for text and document data, audio and video data, and object-oriented data access. Oracle began work on its multimedia server several years ago and leads its competitors in the development of storage and retrieval technologies. This high-throughput, high-speed server technology delivers video on demand to consumers, using set-top, "smart TV" processors as clients.

NOTES

5.5 SQL WINDOWS AND POWER BUILDER

Team Developer is a desktop Rapid Application Development tool, providing both standalone Windows application and Microsoft .NET deployment. Team Developer originally came from SQLWindows, and joined the Unify product family when they acquired Gupta Technologies.

SQLBase

'Embedded Database' Very Low Total Cost of Ownership (TCO) SQLBase effectively reduces the development, deployment and maintenance costs of business applications. The ultra-low TCO of SQLBase ensures that ISVs and end-customers can maximize their ROI on business database applications. SQLBase offers very competitive deployment pricing; including a pay as you go licensing model that allows you to buy the number of licenses needed without any large upfront payments. SQLBase is known for its DBA-less operations that reduce database administration almost to zero. This is an important feature for vertical solutions, as it decreases the cost associated with end-user support as well as making it more cost-effective for large end-user deployments.

This section will introduce you to developing applications with Gupta SQLWindows, client/server application development and deployment environment for Microsoft Windows.

In SQLWindows lets you code and layout applications in an integrated graphical environment:

NOTES

SQLWindows includes the following features:

- Drag-and-drop user interface design.
- Multiple source code and user interface views.
- Context-sensitive coding assistant.
- Online help for all functions.
- Interactive debugging.

Complete the following steps to open, compile, and use the completed Account Info application in SQLWindows:

1. Launch **SQLWindows**. The default Start menu location is Programs, Gupta, Team Developer 3.0, SQLWindows 3.0.
2. Select **File, Open**.
3. Navigate to the **\Gupta\Samples** directory.
4. Double-click **AccountInfo.app**.

The Account Info application opens. This application requests a valid Logon and then allows you to view account information for companies stored in the Island database.

5. Select **Debug, Go**.

The application compiles, and you view the Database login dialog.

Power builder

PowerBuilder is an integrated development environment owned by Sybase, a division of SAP. It has been in use since 1991, peaking around 1998 with around 100,000 users.

While Powerbuilder's market share has diminished, many applications created with it are still in use. In 2010, Sybase released a major upgrade to PowerBuilder, intended to compete directly with Microsoft Visual Studio.

Sybase sells another programming language called PocketBuilder. It is based on PowerBuilder and used for creating applications that run on mobile devices such as cell phones or PDAs.

On April 20, 2010, Sybase released PowerBuilder version 12.0, a major upgrade. It includes the ability to create applications that target several different platforms offered by the .NET Framework, including ASP.NET Web Forms, Windows Forms, or WPF. These features are offered through the new PB.NET IDE, which is built on the Visual Studio Isolated Shell. This latest release is intended to better align PowerBuilder with its competitors while continuing to offer its most distinguishing product: the DataWindow. In PowerBuilder 12, PowerScript is a fully CLS-compliant .NET language and comes with the ability to deploy fully managed code. The older PowerBuilder IDE (now termed "PowerBuilder Classic" by Sybase) continues to be made available for vanilla PowerBuilder development.

PowerBuilder has a native data-handling object called a **DataWindow**, which can be used to create, edit, and display data from the database. This object gives the programmer a number of tools for specifying and controlling

user interface appearance and behavior, and also provides simplified access to database content. To some extent, the DataWindow frees the programmer from considering the differences between Database Management Systems from different vendors.

PowerBuilder also includes a scripting language, PowerScript, which is used to specify the application behavior when events occur. Events usually correspond to user actions, such as clicking on an element of the user interface, or closing a window. PowerBuilder 10 and above use Unicode internally. This allows the use of characters from multiple languages concurrently. This affects the coding of API calls, however, the upgrade from PowerBuilder 9 to 10 automatically converts your existing API calls for you. Therefore, in PB10, you have a choice of making ANSI or Unicode API calls at the call level. PowerBuilder 10 supports reading and writing to flat files larger than 2 GB. PB8.0 supports the TRY CATCH FINALLY syntax of error handling. PB 10 has a built-in date-picker control.

Main features of PowerBuilder

1. PowerBuilder is a rapid 4GL integrated application development tool which you can use to develop wide range of business applications from Windows to Web development. 3GL Coding and Embedding SQL is the best feature of PowerBuilder development.
2. PowerBuilder.NET IDE is a CLS compliant .net ready and you can develop .net applications that run on any platform that has microsoft .net framework. PowerBuilder ships with two versions for developing Win32 as well as .net applications.
3. PowerBuilder IDE is designed from C/C++ Language and its compiler is very efficient to produce highly optimized code to run on different platforms. Only PowerBuilder has a special feature to organize your objects inside Libraries and compile as a Dynamic Linking Loading (DLL) libraries.
4. PowerBuilder is a cross platform development tool and you can run your application in any platform like Windows, Unix & Linux flavours.
5. PowerScript is the powerbuilder development language which is easy to learn like Visual Basic with the power of C/C++. Powerscript is not case-sensitive and PBNI enable you to talk to external language objects like C++ or Java. PowerBuilder API can call CORBA objects.
6. Powerbuilder supports Object Oriented Programming (OOPD) and confirms Encapsulation, Inheritance, Polymorphism, Overloading and Overriding etc.. You can quickly inherit any visual controls, forms and custom objects and effectively minimize programming code.

NOTES

NOTES

7. PowerBuilder has built in exception handling and supports locale development.
8. Powerbuilder has patented powerful data grid called "Datawindow" which you can use to develop data entry forms, reports, graphs and tree views at a fraction of cost. Also you can create mail merge and rich text editors.
9. PowerBuilding allows you to create native Custom Component development including Visual, Non-Visual and COM/COM+ components. You can also inherit from any controls or from any Microsoft OLE controls.
10. Powerbuilder has all the controls that you want to develop an application. You do not need to run behind third party controls and looking for licensing issues.
11. Easily develop and deploy any Client/Server or 3 Tier applications. It reduces development time 40% less than Visual Studio and 60% less than any java development.
12. Powerbuilder connects to almost all types of databases including Sybase, Microsoft SQL, Oracle, Informix, MySQL and SQL Anyware through PB Native, Direct Connect, Microsoft OLE, SQL.NET, ODBC, and JDBC drivers. It also supports connecting to web services.
13. PowerBuilder is 100% Enterprise ready. You can develop any 3 Tier distributed applications using powerful EAServer application server.
14. PowerBuilder has a unique "Data Pipeline" object which you can use to meet all the ETL (Extract, Transform and Load) requirements without requiring any third party softwares.
15. PowerBuilder IDE has all the features that are available in Microsoft Visual Studio including intellisense.
16. Powerbuilder.Net allows you to create Win forms, Datawindow.Net, WPF and WCF Services. You can apply C# object-oriented programming techniques such as interfaces, delegates, generics etc.,
17. Pocketbuilder, a short cut version of powerbuilder to develop mobile applications.

PowerBuilder is used primarily for building business applications. PowerBuilder is used by some companies in the financial and telecoms sectors. However in recent years, PowerBuilder has seen substantial competition arise from Java development tools and from Visual Studio.

There are a number of third-party tools that build upon and enhance the functionality of PowerBuilder, including Appoon for PowerBuilder, Visual Expert for PowerBuilder, and Enable Multilingual.